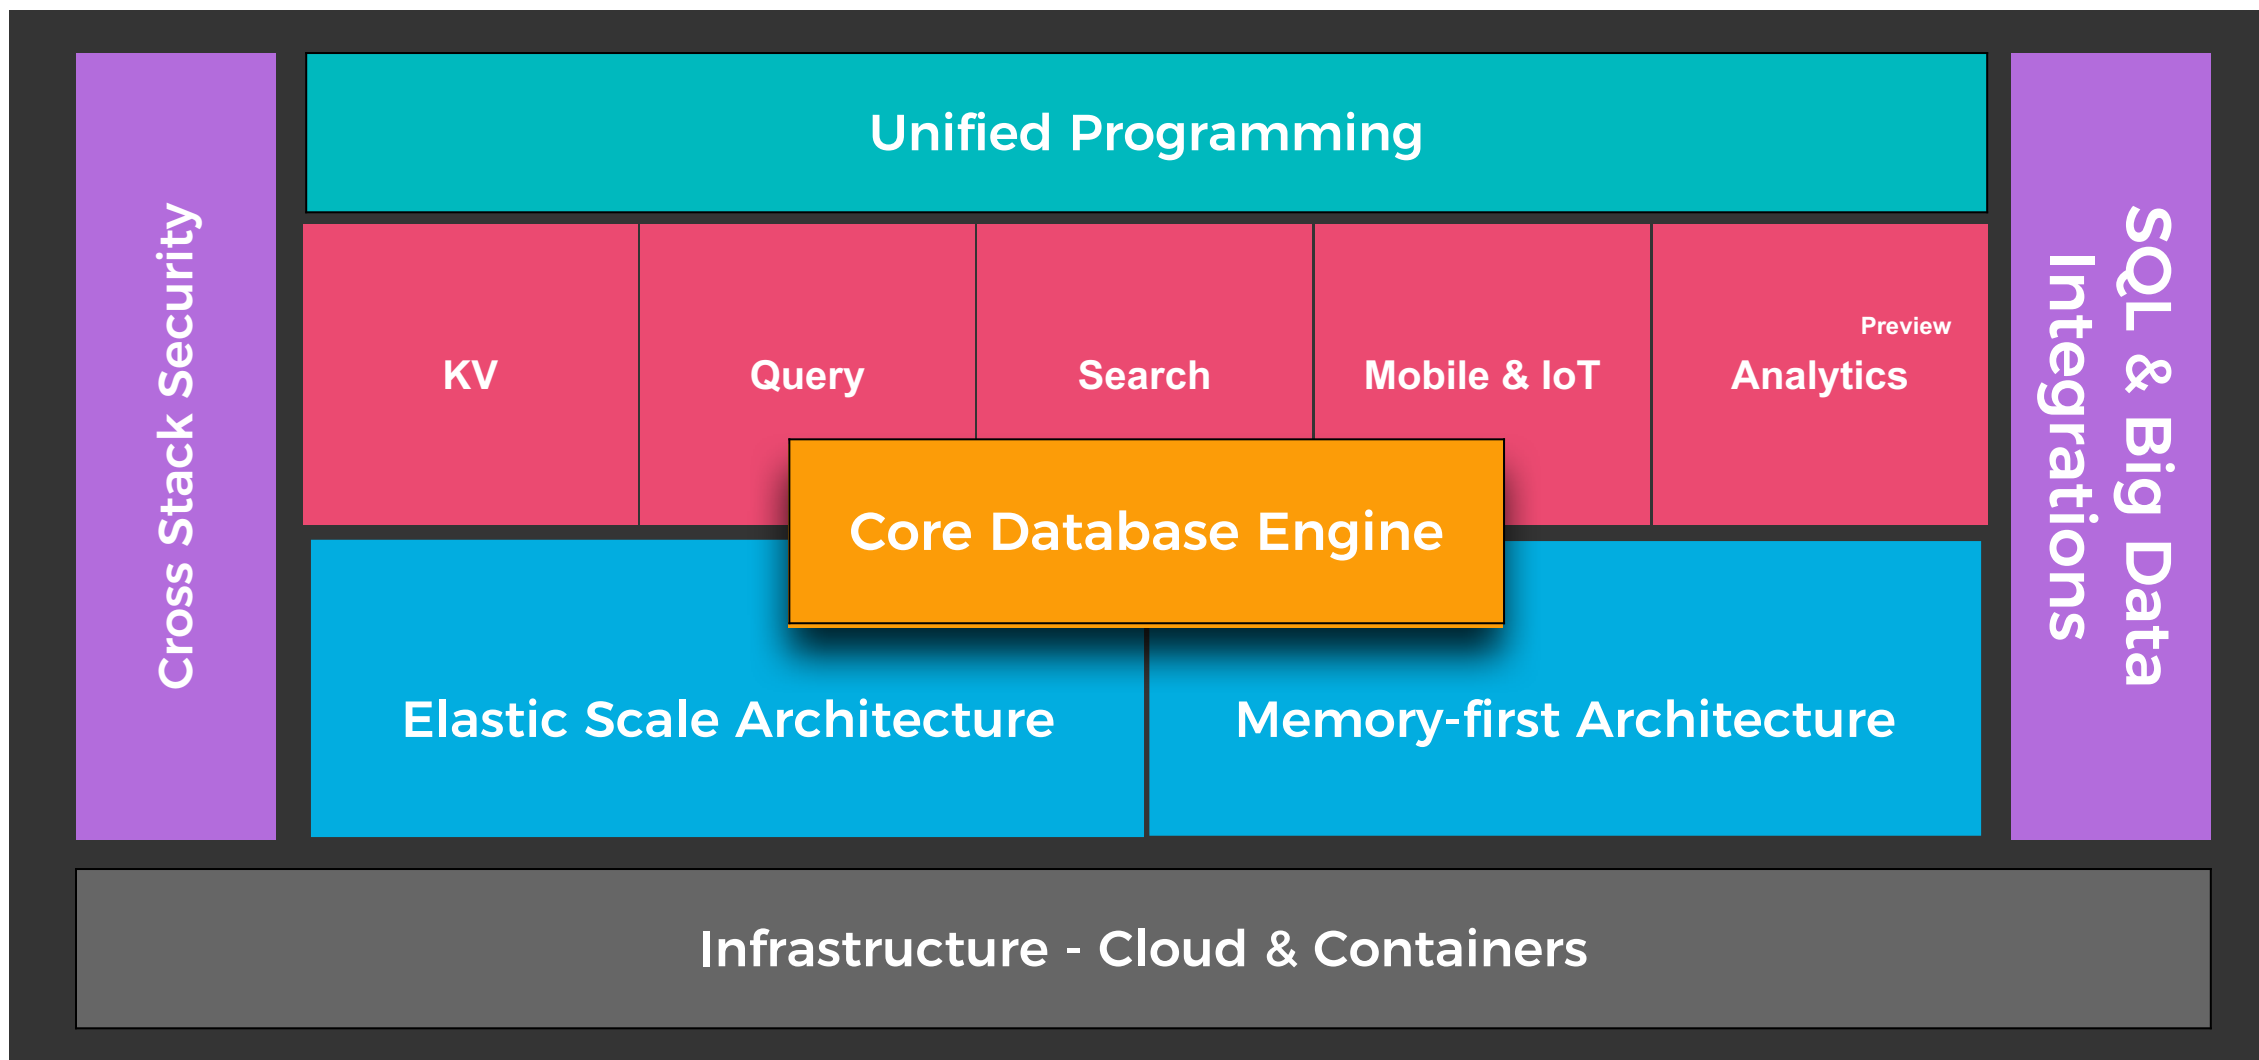


Architecture and Administration Basics

Workshop Day 1 - Architecture

Couchbase Data Platform

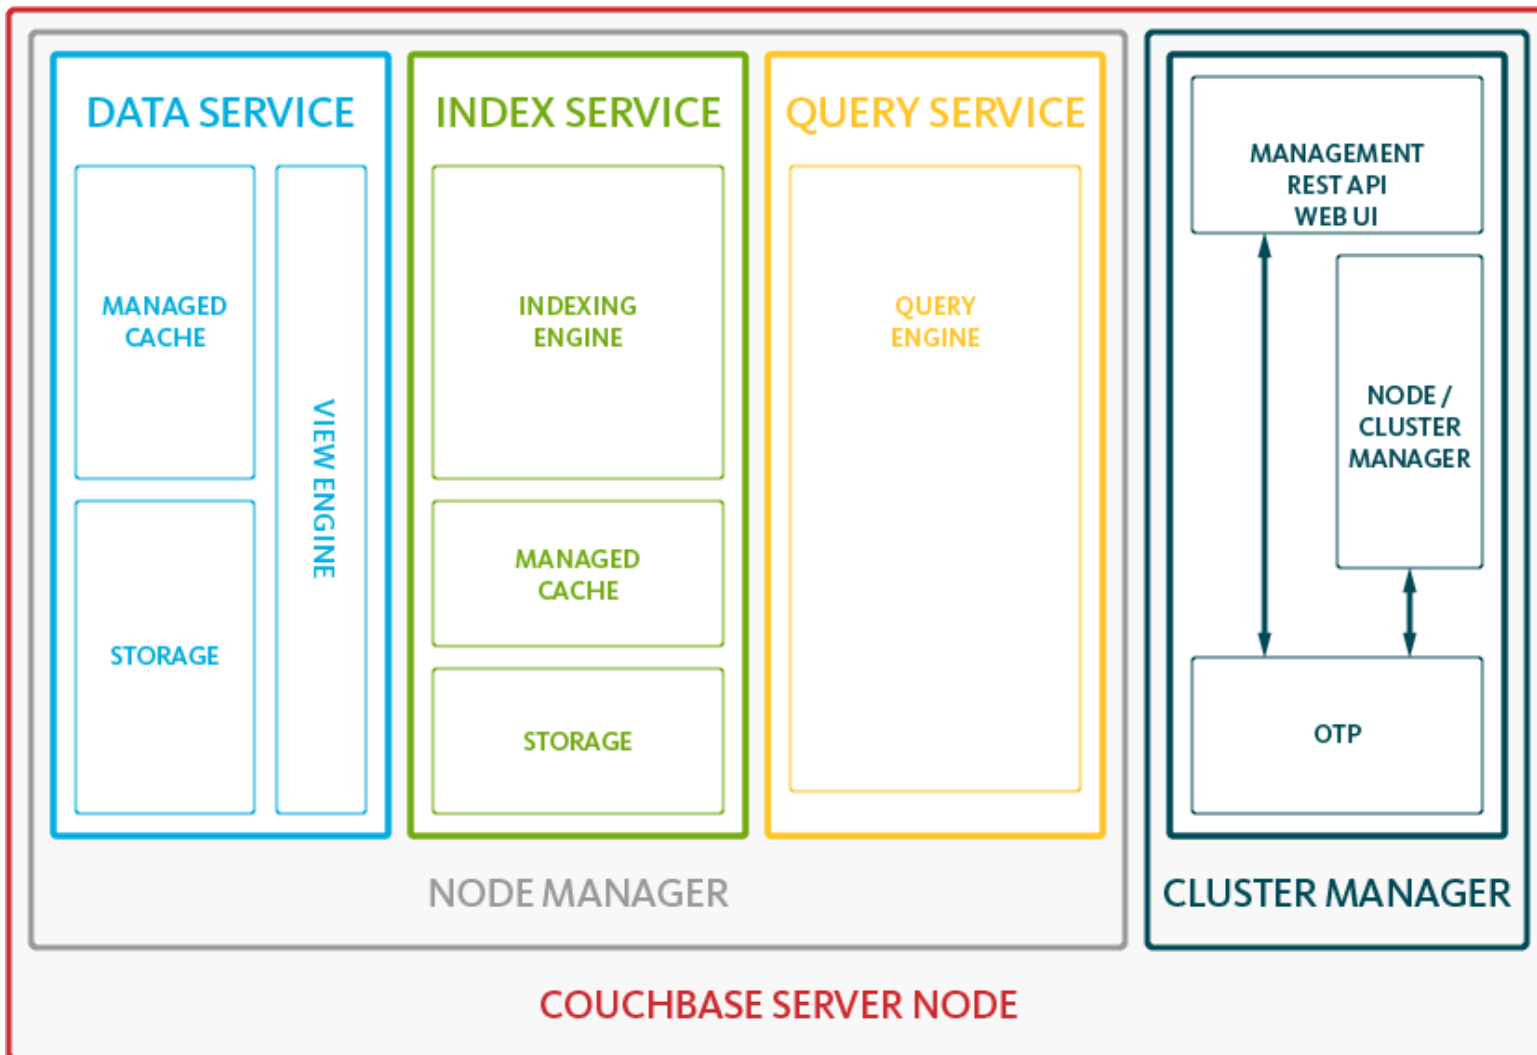




1


Architecture

Couchbase Architecture



- **Data Service** – Key Value Store and builds and maintains Distributed secondary indexes (MapReduce Views)
- **Indexing Engine** – builds and maintains Global Secondary Indexes
- **Query Engine** – plans, coordinates, and executes queries against either Global or Distributed indexes
- **Cluster Manager** – configuration, heartbeat, statistics, RESTful Management interface

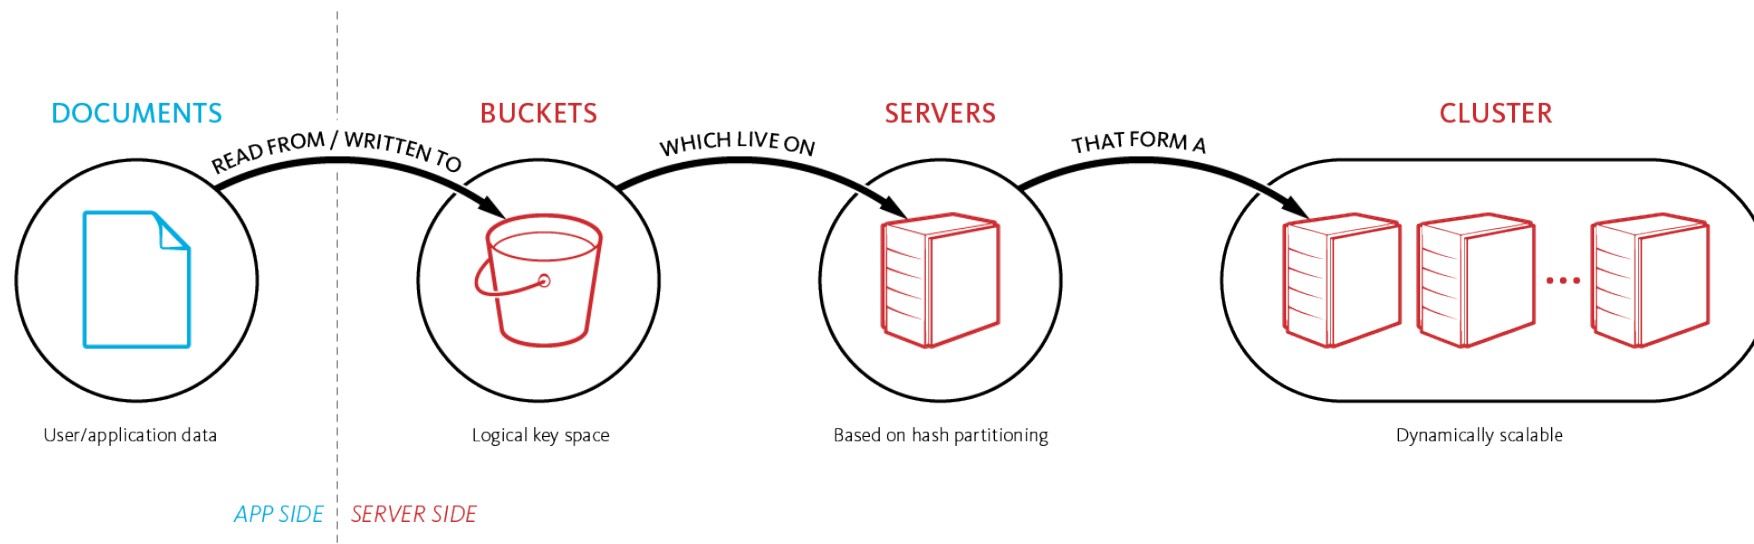
Storing And Retrieving Documents






Buckets - When to use more than one ?

- When you need to treat or access the data differently
 - Different High Availability requirements (1,2 or 3 replicas)
 - Different performance / residency needs (how much data to cache)
 - Security / Multi-tenancy
 - Segregating Binary and JSON data – especially with view usage



Auto sharding - Bucket and vBuckets



Virtual buckets

- Bucket
 - A **bucket** is a logical, unique key space
 - Multiple buckets can exist within a single cluster of nodes
- vBuckets
 - Each bucket has active and replica data sets (1, 2 or 3 extra copies)
 - Each data set has **1024 Virtual Buckets** (vBuckets)
 - Each vBucket contains 1/1024th portion of the data set
 - vBuckets do not have a fixed physical server location
 - Mapping between the vBuckets and physical servers is called the **cluster map**
 - Document IDs (keys) always get hashed to the same vBucket (consistent hashing)
 - Couchbase SDK's lookup the vBucket -> server mapping



Bucket Comparison

	Memcached	Couchbase	Ephemeral <small>New in 5.0</small>
Persistence	X	✓	X
Replication	X	✓	✓
Rebalance	X	✓	✓
XDCR	X	✓	✓
N1QL	X	✓	✓
Indexing	X	✓	✓*
Max Object Size	1MB	20MB	20MB

* MOI, FTS only



Ephemeral Bucket Benefits and Limitations

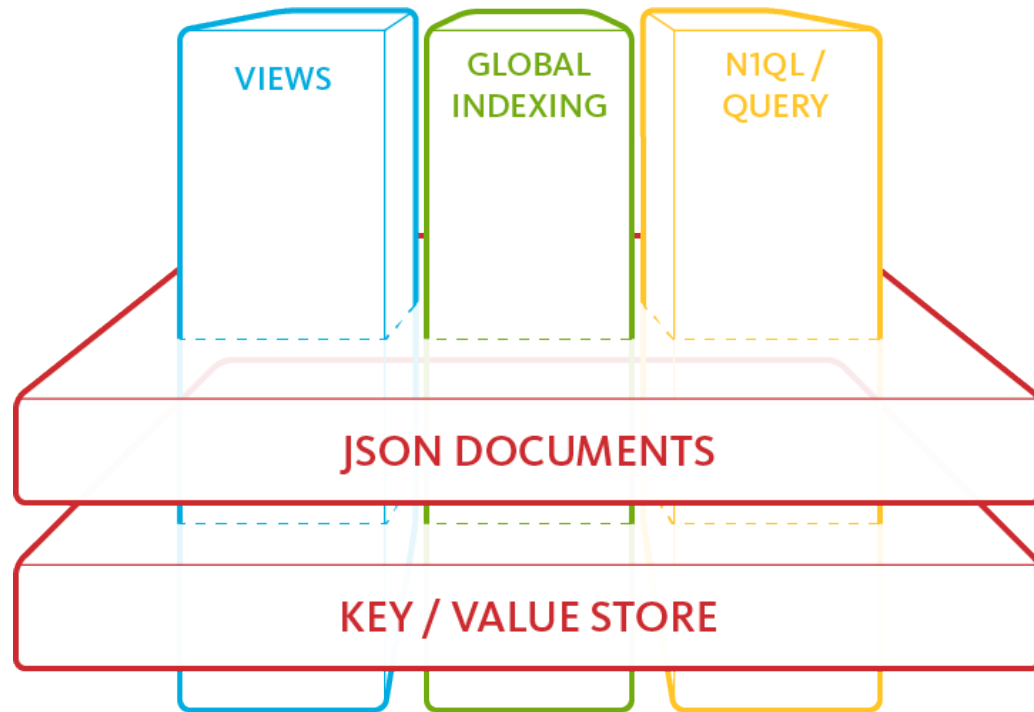
- **Benefits**

- No high performance disk subsystem required
 - Lower cost VMs
 - Smaller chassis
- Even more consistent high performance
 - No disk IO contention (i.e. compaction)
- Lower CPU consumption
 - No Disk Write Queue
 - No IO threads
- Faster maintenance operations
 - No warm-up
 - Faster node restart
 - Faster rebalance – currently 4x faster in our lab!

- **Limitations**

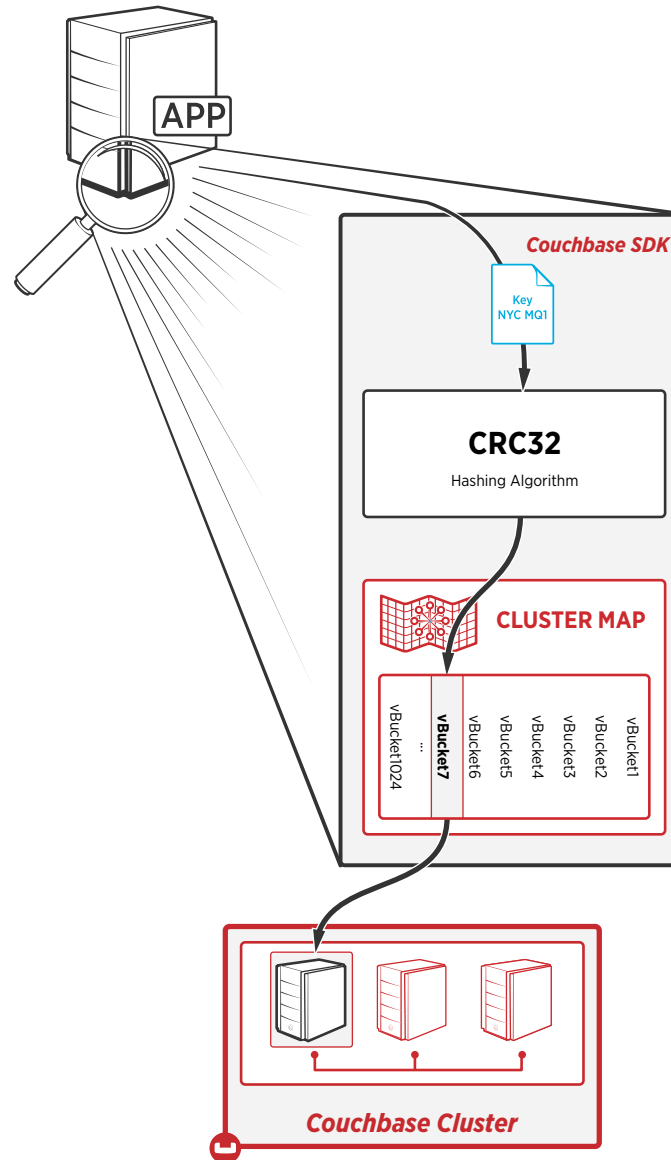
- Data set must fit in memory
 - Configurable OOM handling
- No automatic recovery from total power loss
 - Backups and XDCR still supported!
- XDCR limitations
 - Ephemeral to Ephemeral
 - Couchbase to Ephemeral
 - Ephemeral to Couchbase NOT supported
- Only Memory Optimized Indexes (MOI) and Full Text Search (FTS) are supported
 - No Views or Standard GSI

Couchbase Data Access

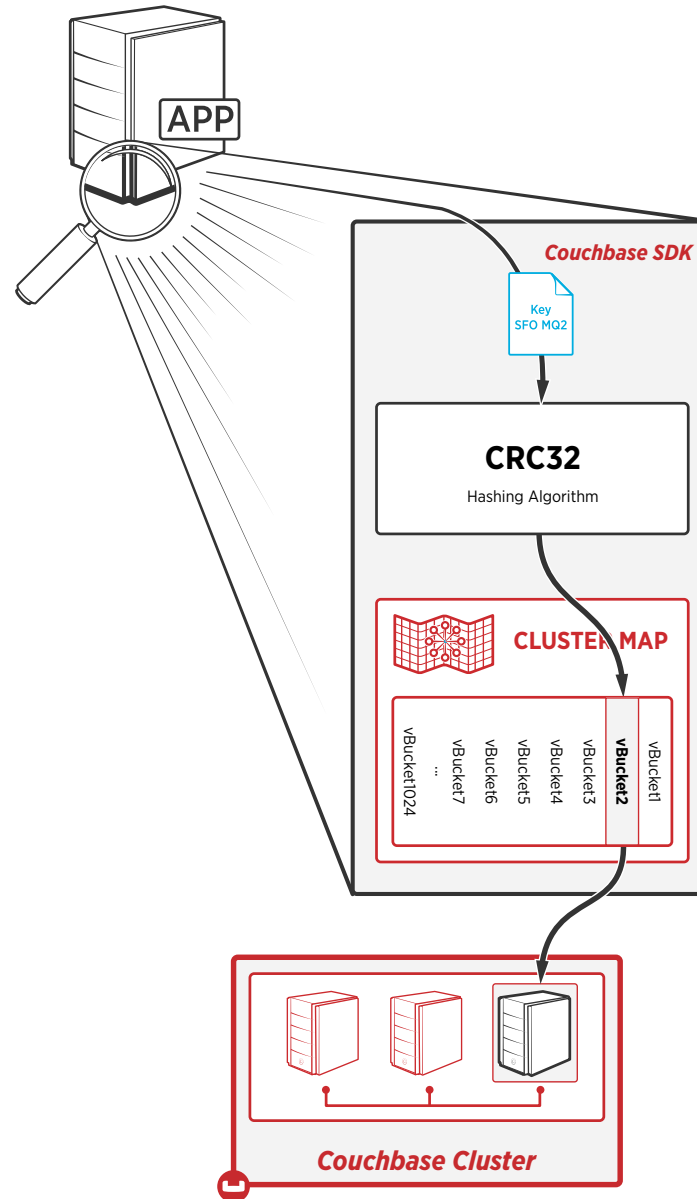


- Everything is layered on top of **Key Value**
- A **Document** store is a special case of Key-Value
- **Views** provide aggregation and real-time analytics through incremental map-reduce
- **Global Secondary Indexes** provide low latency/high throughput indexes
- **N1QL** is a language that provides a powerful and expressive way of accessing documents

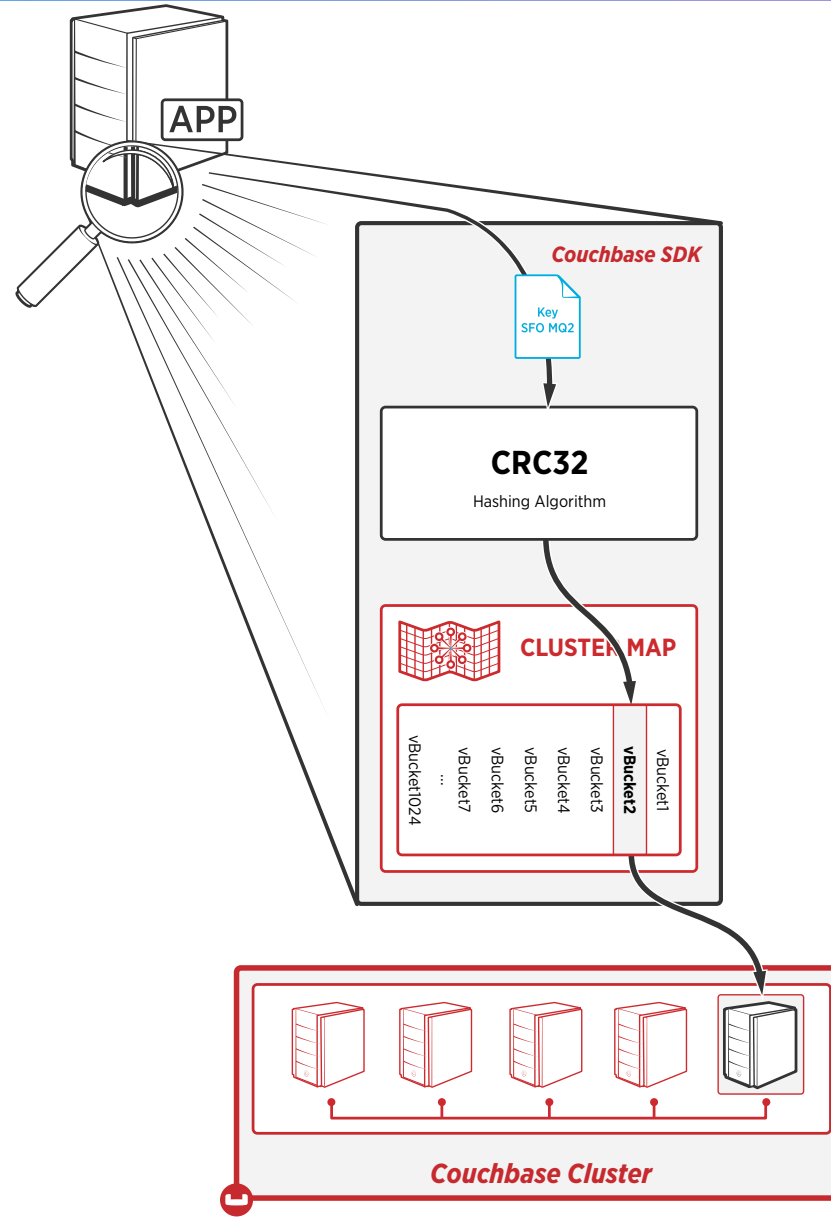
Cluster Map




Cluster Map



Cluster Map - Addition of 2 Nodes



Application to Database Interaction







2

SDK Operations


Couchbase Read Operation



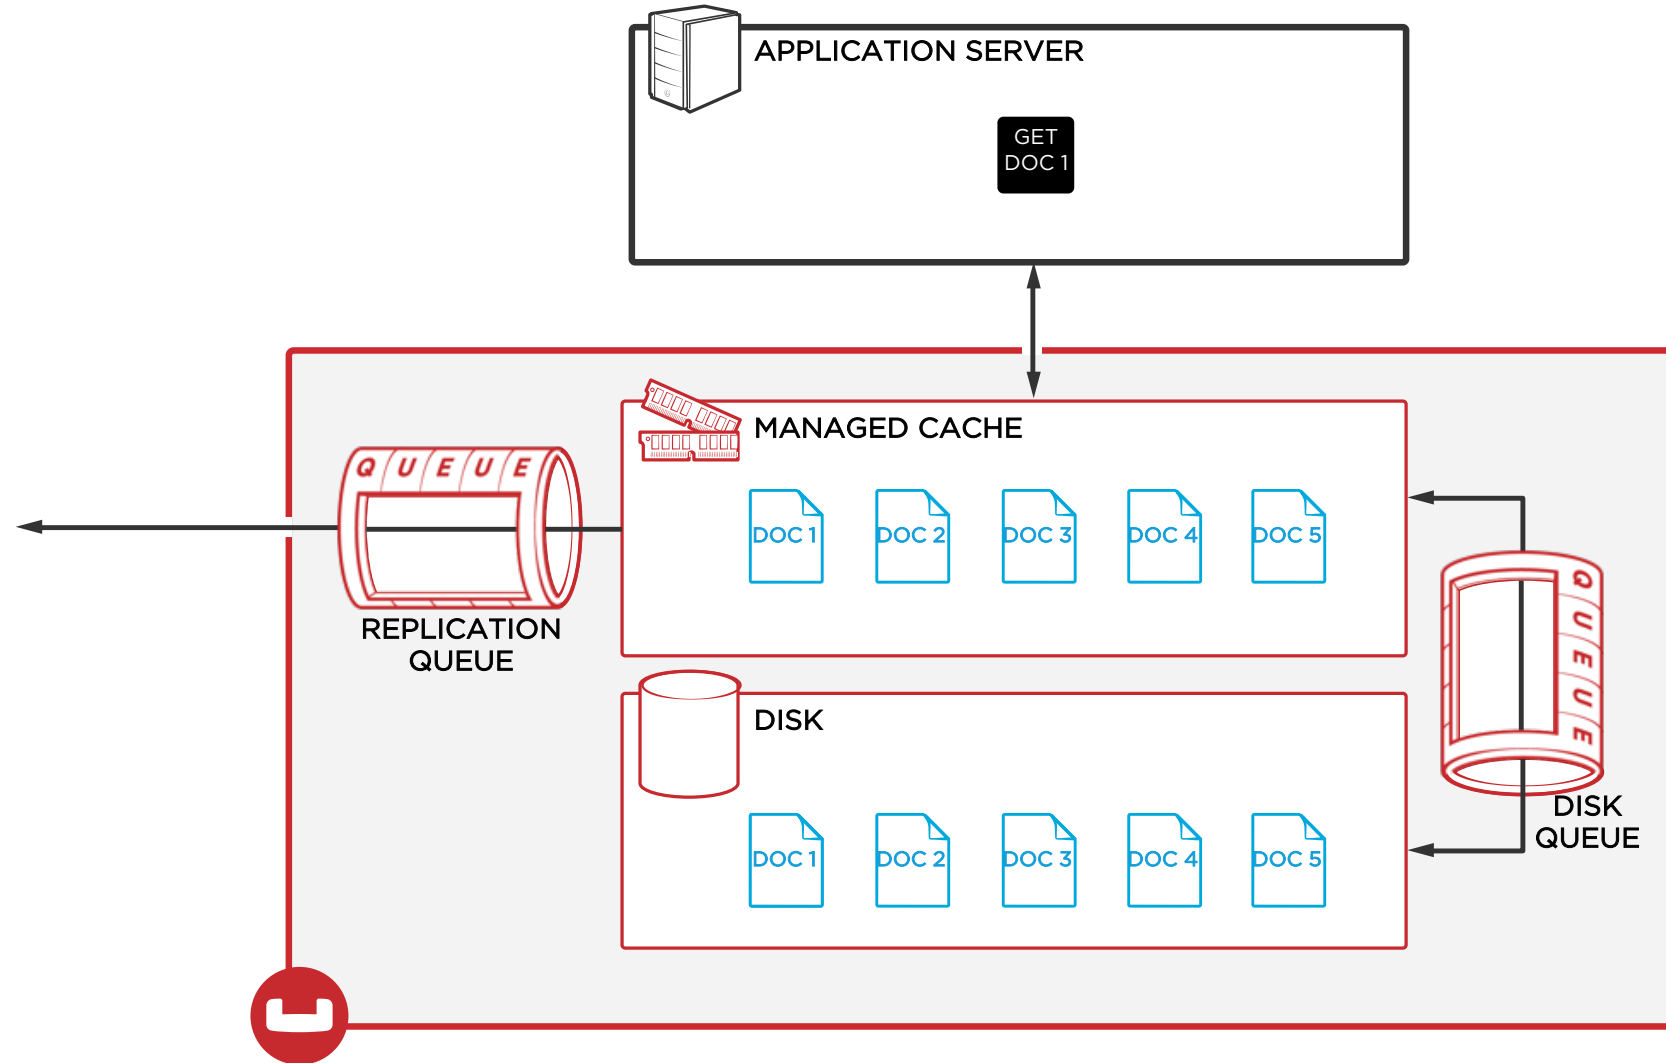
Write Operation



Cache Ejection



Cache Miss

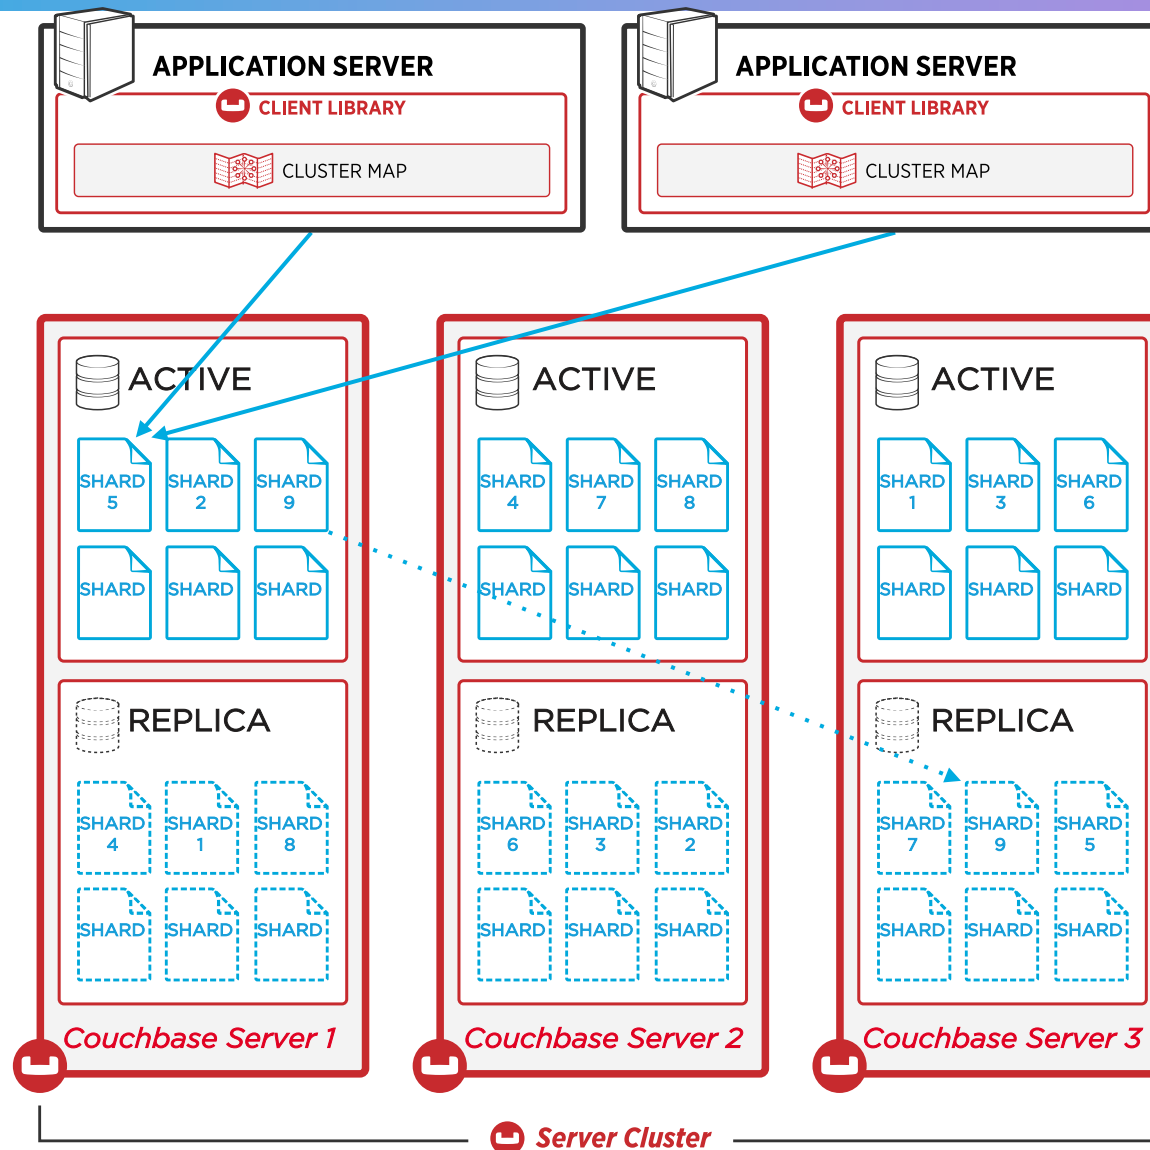




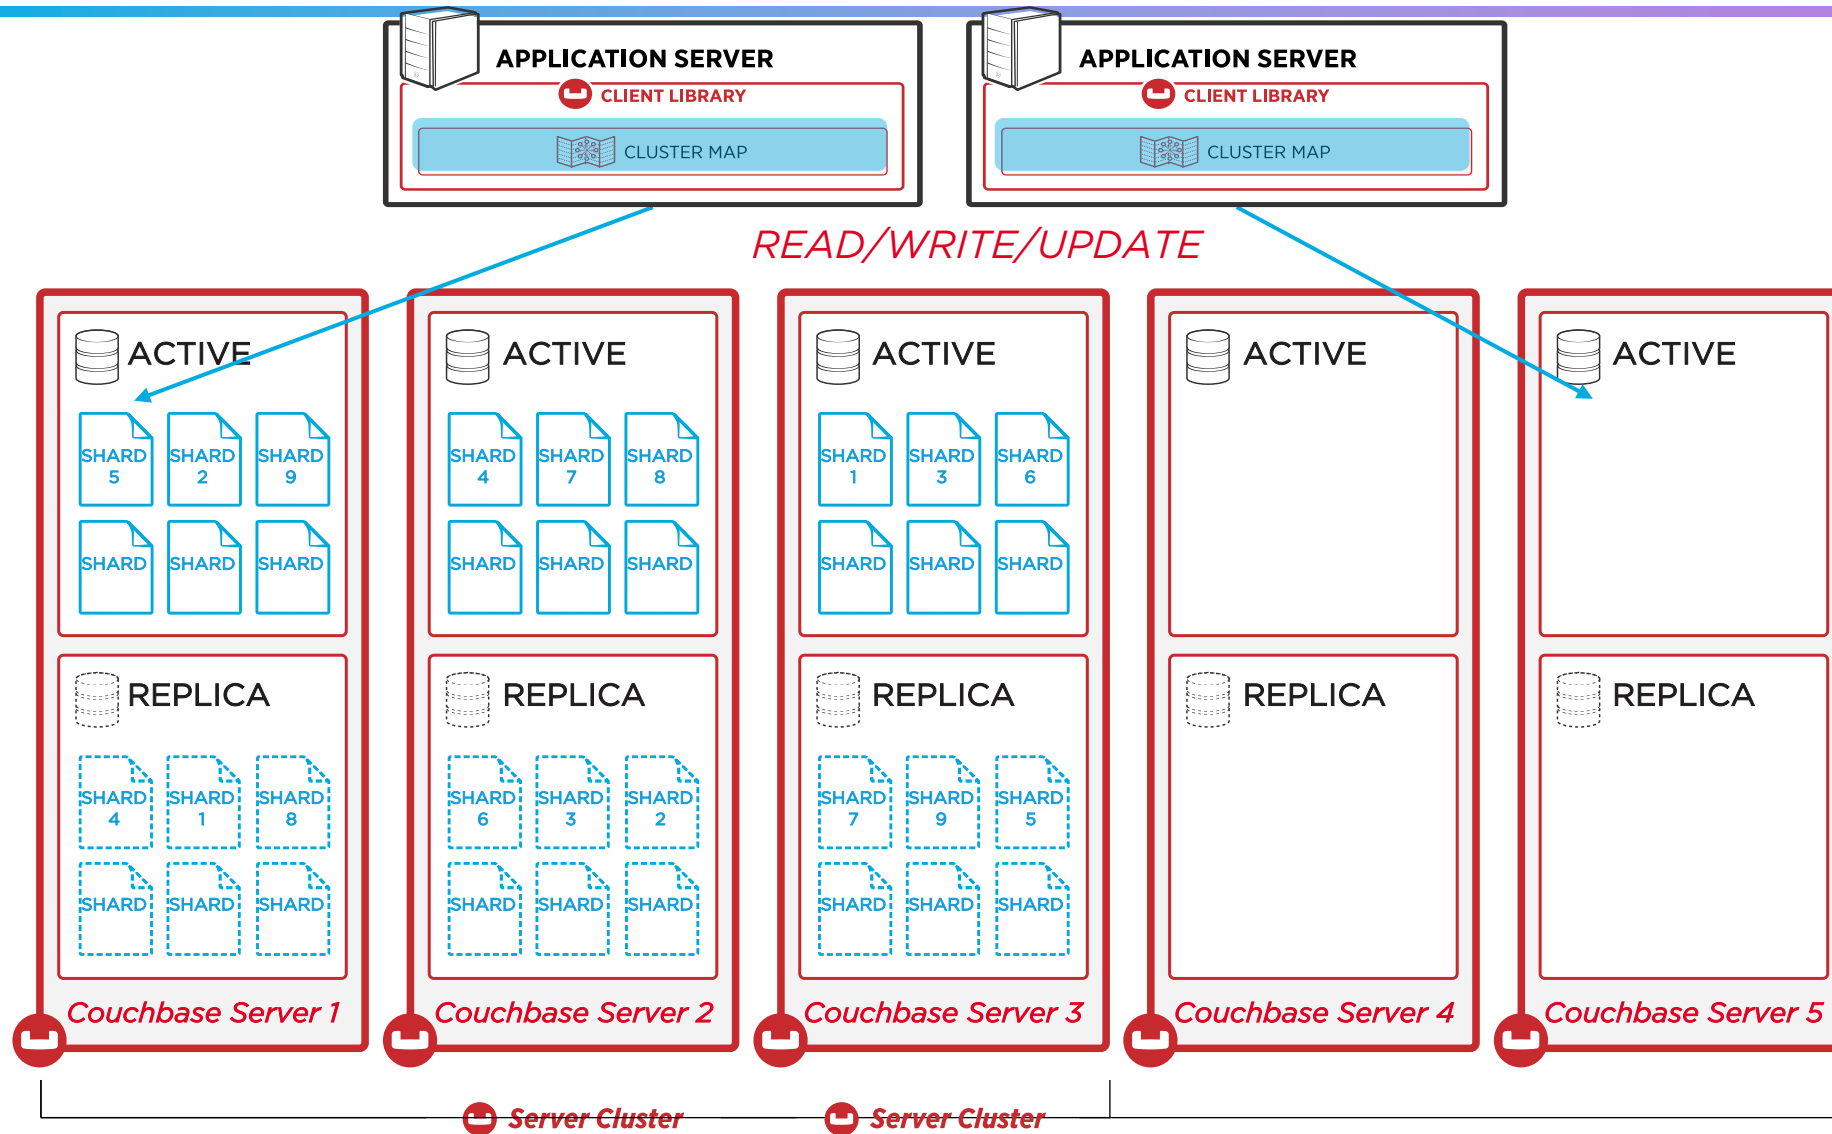
3

Cluster Operations

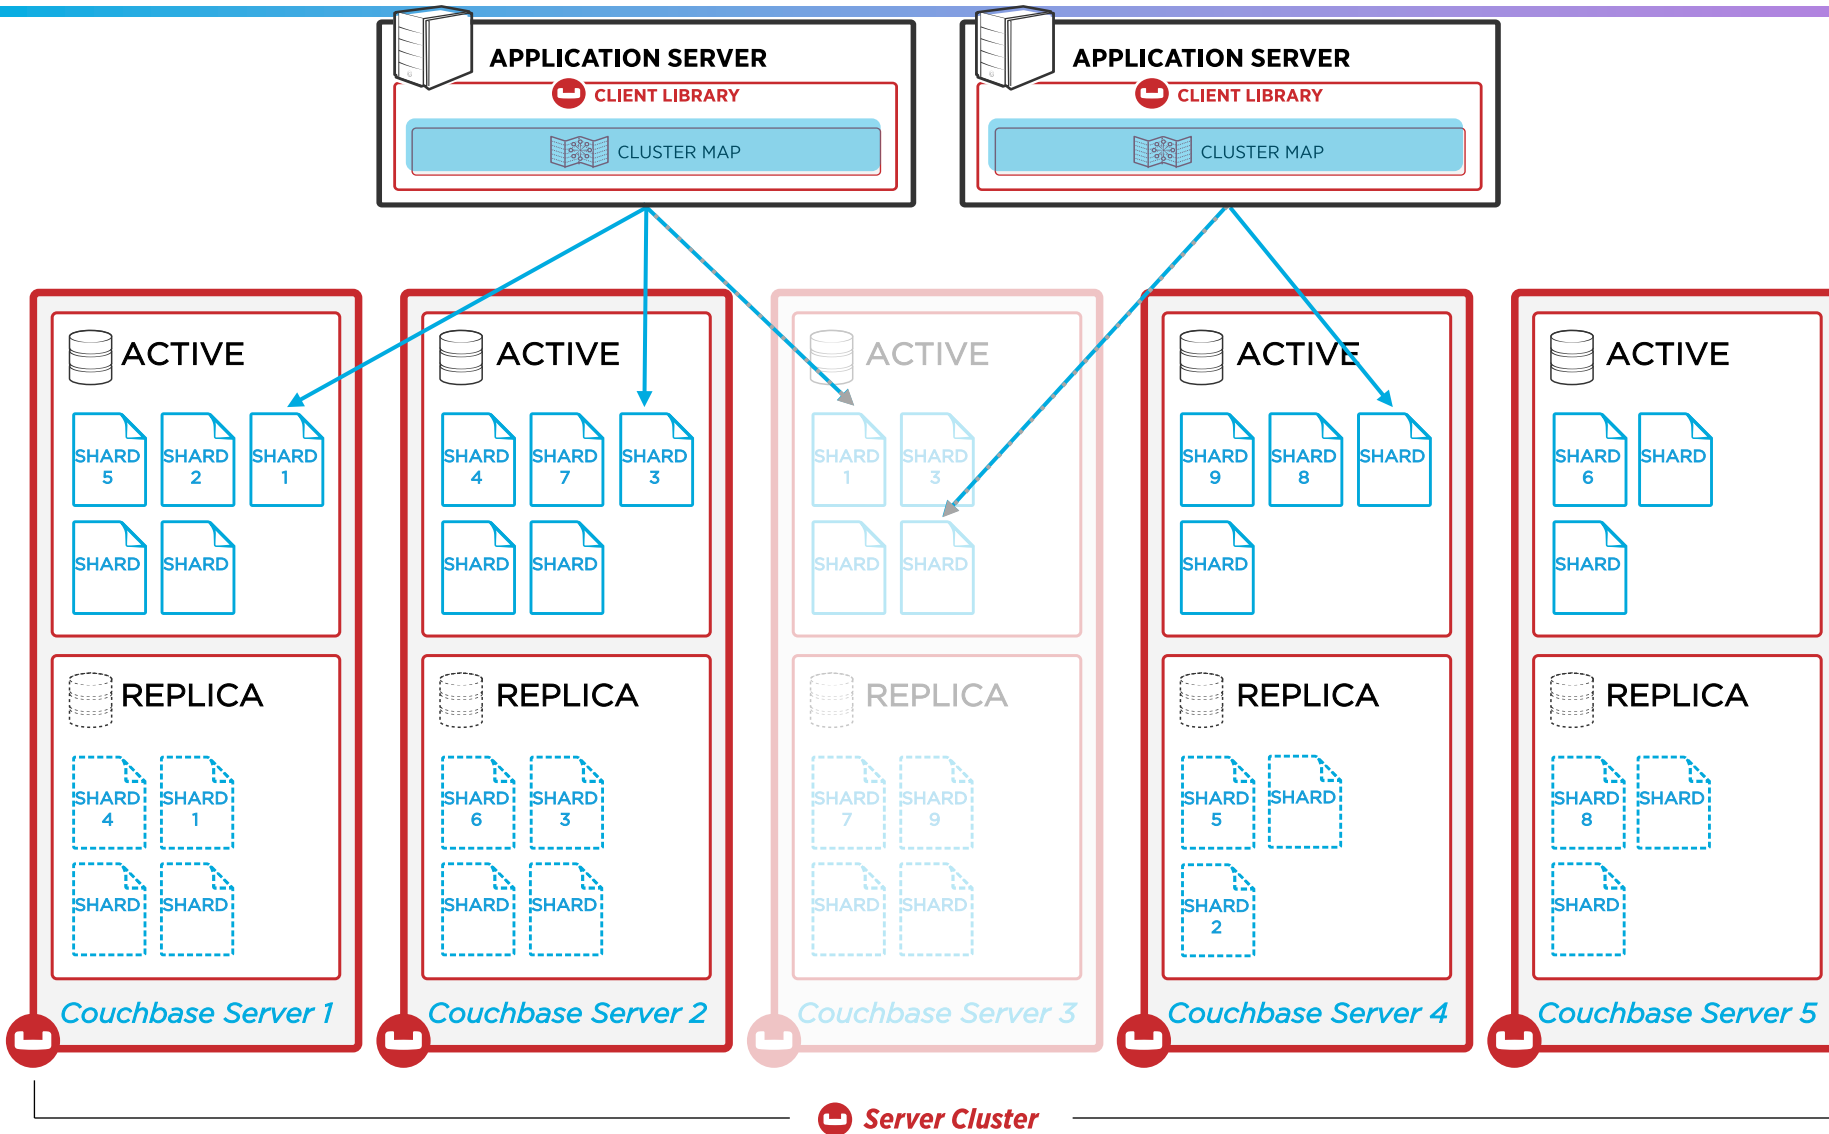
Basic Operation



Add Nodes to Cluster



Fail Over Node





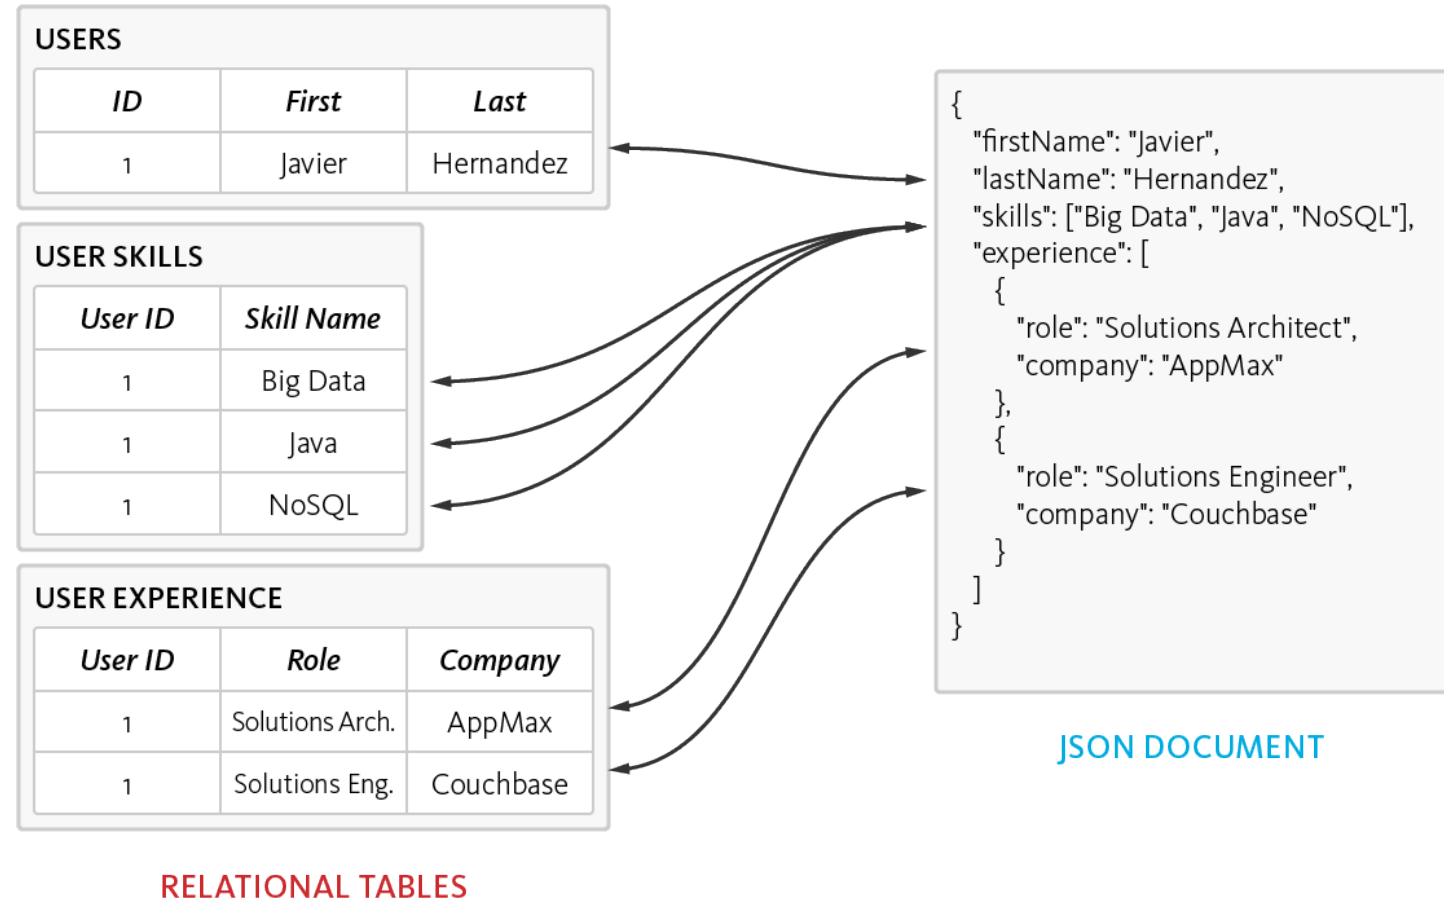
JSON Support

Flexibility

The Power of the Flexible JSON Schema



- Ability to store data in multiple ways
 - Denormalized single document, as opposed to normalizing data across multiple table
 - Dynamic Schema to add new values when needed



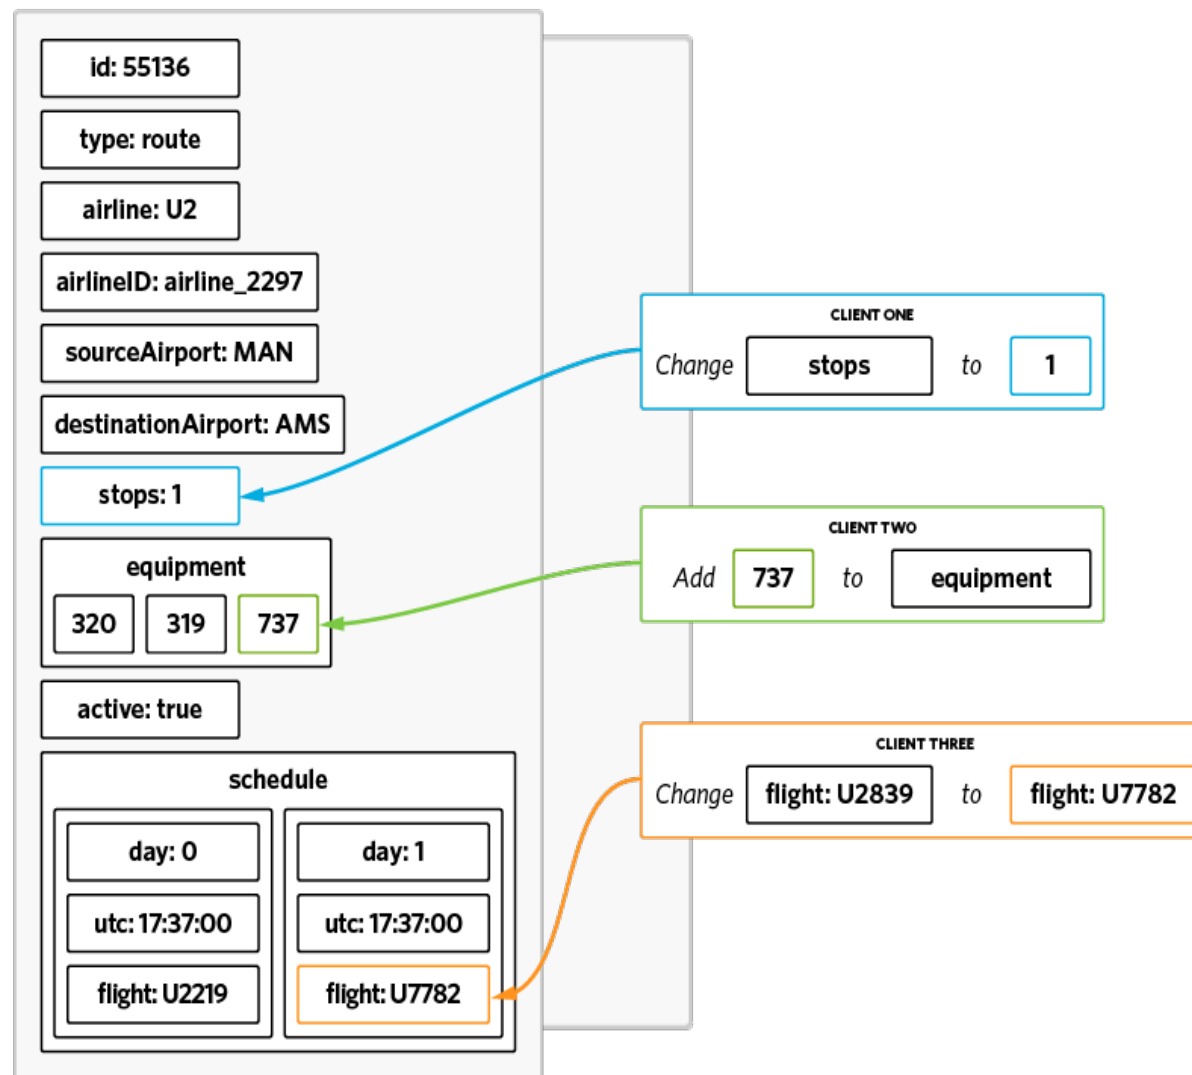
Efficient Sub-Document Operations



- Document Mutations:

- Atomic Operate on individual fields
- Identical syntax behavior to regular bucket methods (upsert, insert, get, replace)
- Support for JSON fragments.
- Support for Arrays with uniqueness guarantees and ordinal placement (front/back)


Document Key: "route_55136"



Faster Key Based Operations!



10-14X Faster Document Read & Update Operation





Nickel (N1QL) : SQL-Like Querying Support

- SQL-like Query Language
 - Expressive, familiar, and feature-rich language for querying, transforming, and manipulating JSON data
 - ANSI 92 SQL Compatible – Selects, Inserts, Updates, Group By, Sort, Functions etc.
- N1QL extends SQL to handle data that is:
 - **Nested**: Contains nested objects, arrays
 - **Heterogeneous**: Schema-optional, non-uniform
 - **Distributed**: Partitioned across a cluster

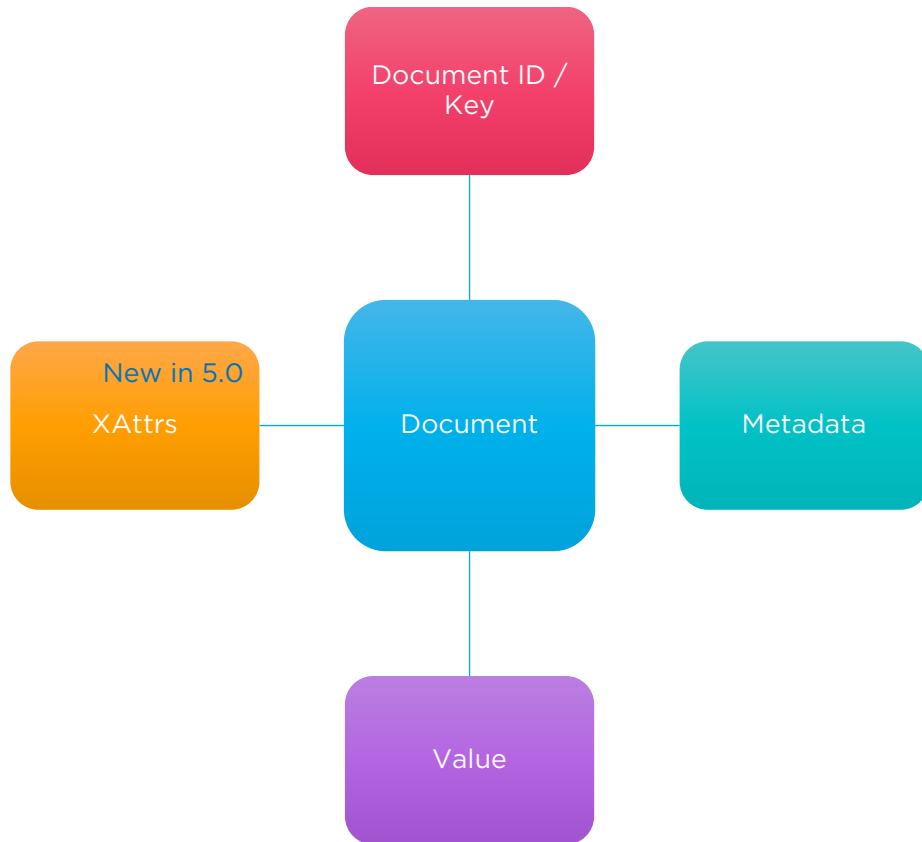


**Power of
SQL**



**Flexibility
of JSON**

JSON Document Support



- **Document ID / Key (Max 250 bytes):**

- Must be unique / Lookup is extremely fast
- Similar to primary keys in relational databases
- Documents are partitioned based on the document ID

- **Value (Max 20 MB)**

- JSON
- Binary - integers, strings, booleans
- Common binary values include serialized objects, compressed XML, compressed text, encrypted values

- **Metadata (Fixed 56 bytes)**

- CAS Value (unique identifier for concurrency)
- TTL
- Flags (optional client library metadata)
- Revision ID #

- **XAttr (Max 20 MB)**

- Non-enumerable eXtended Attributes

New in 5.0

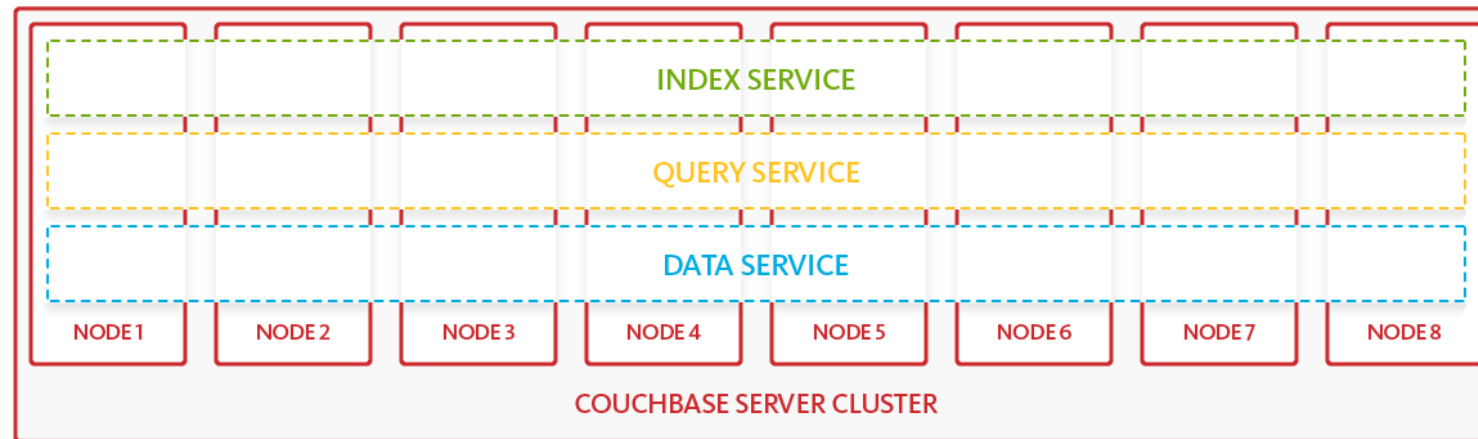


4 | MDS

Modern Architecture – Multi-Dimensional Scaling



MDS is the architecture that enables independent scaling of data, query, and indexing workloads while being managed as one cluster.



Modern Architecture – Multi-Dimensional Scaling




Independent Scalability for Best Computational Capacity — *per Service*.

Heavier Indexing?

Scale up or out
Index Service Nodes.

More RAM for Query Processing?

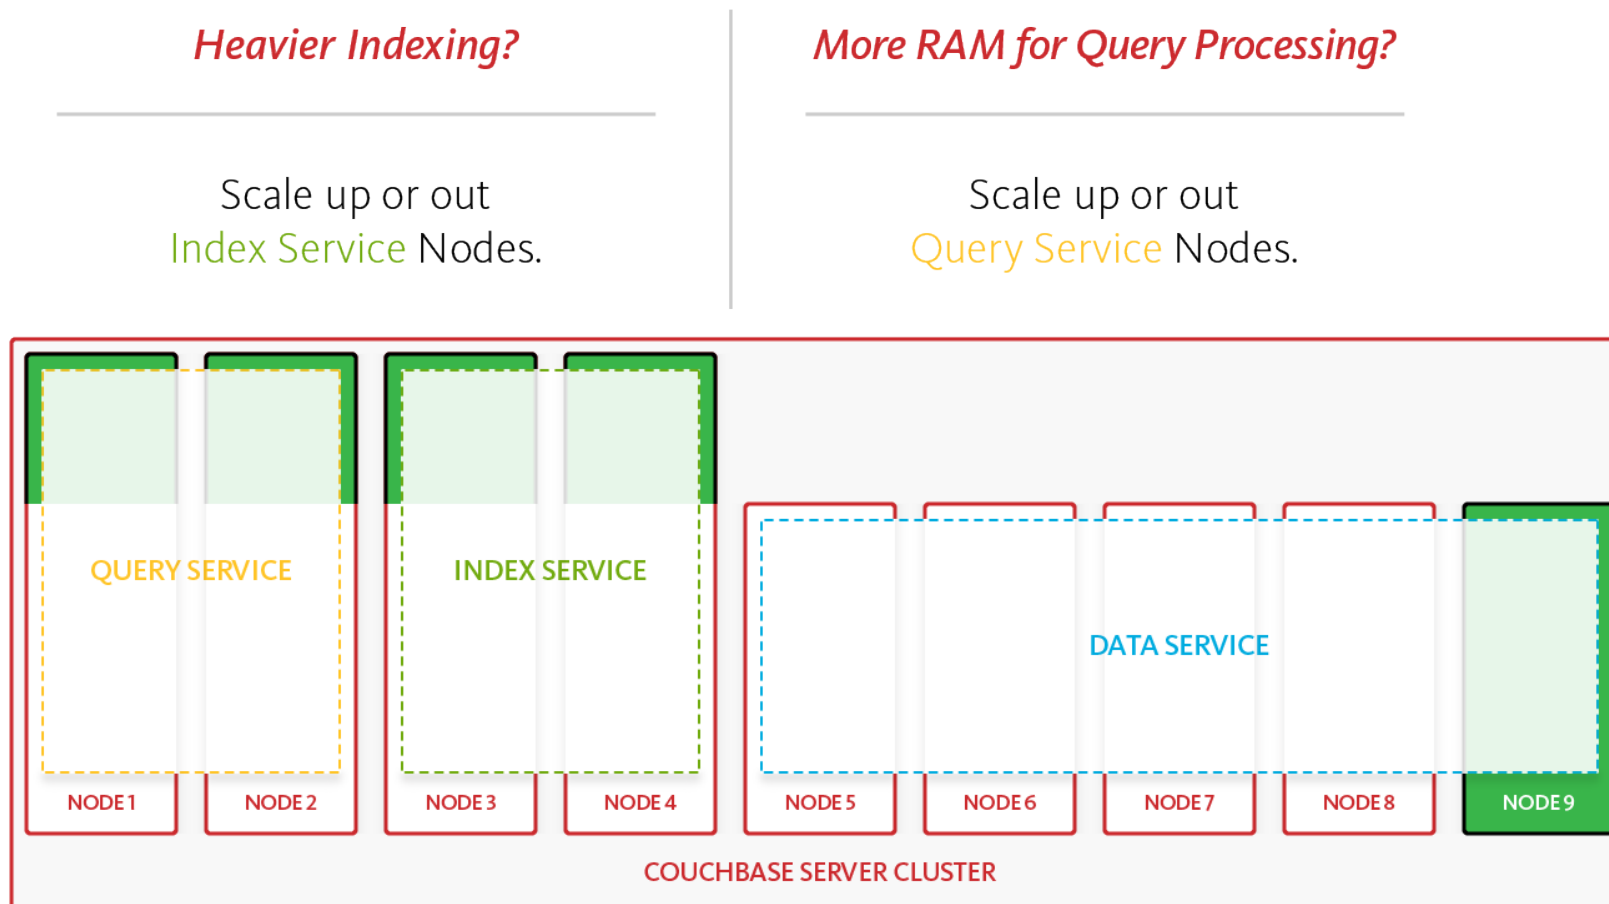
Scale up or out
Query Service Nodes.



Modern Architecture – Multi-Dimensional Scaling



Independent Scalability for Best Computational Capacity — *per Service*.



The same applies for other services




5

Core Principles

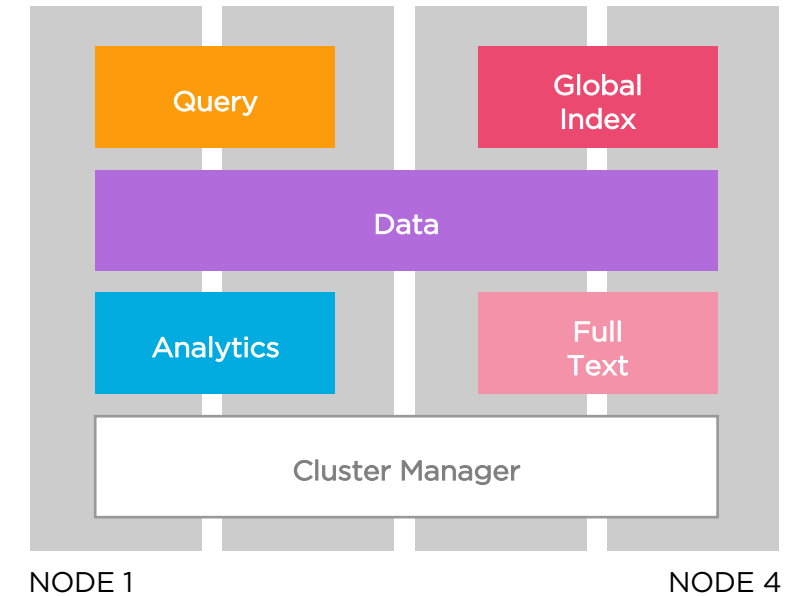
#1 Elastic Scaling Architecture



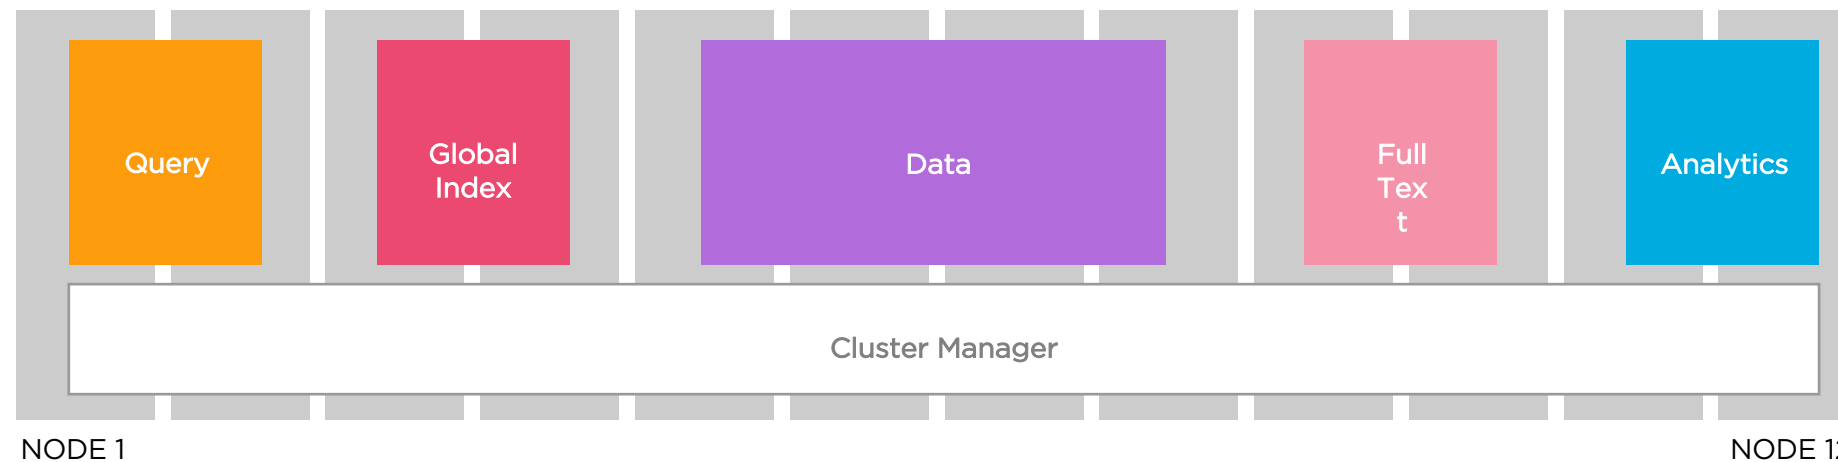
Sample Dev Setup



Sample QA Setup



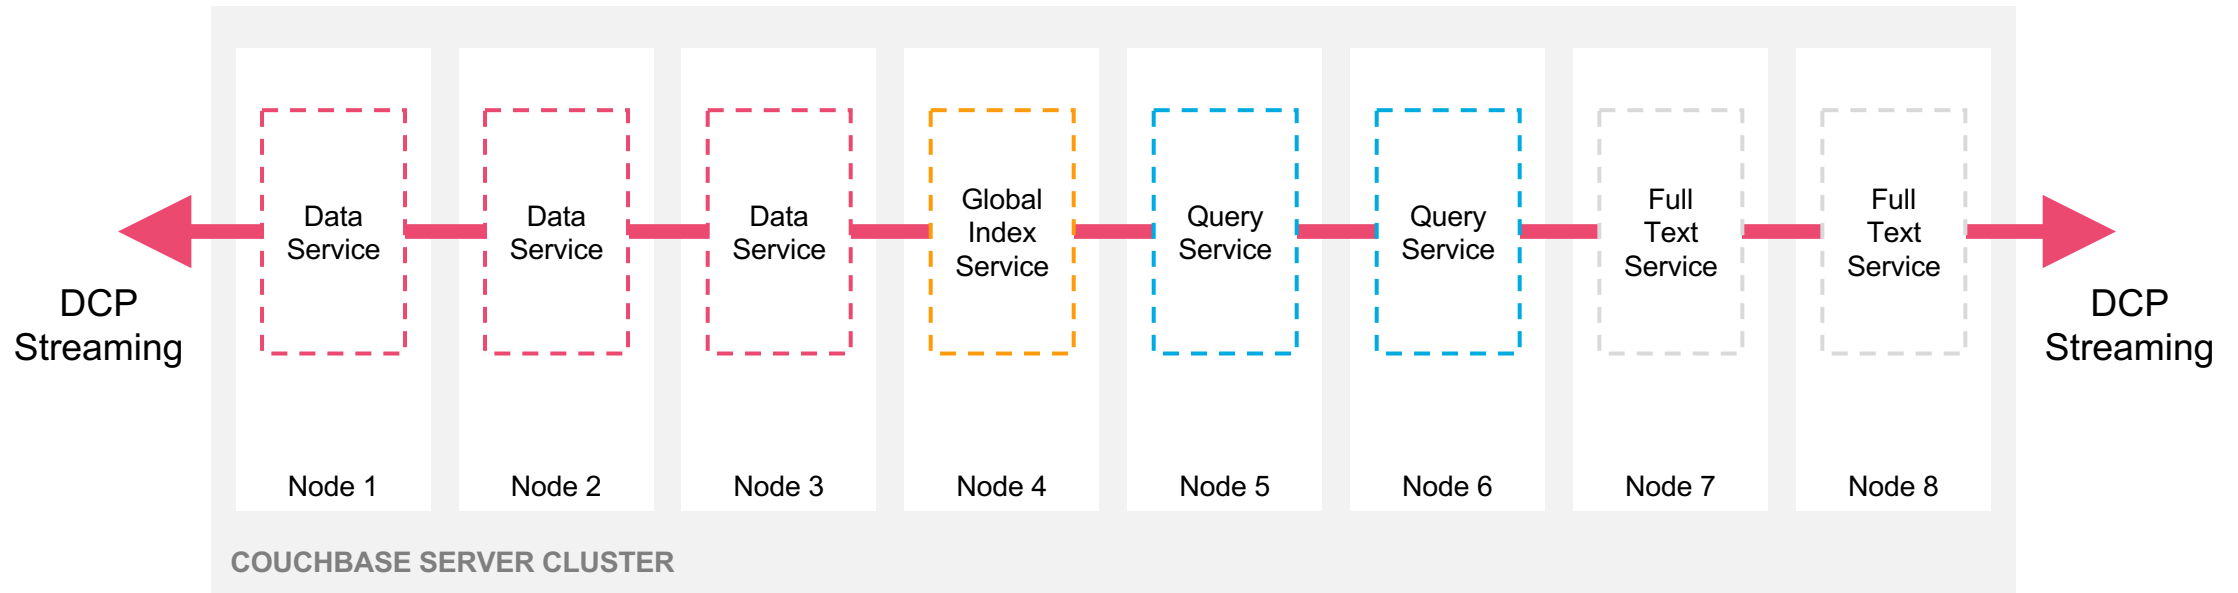
Sample Production Deployment



#2 Memory-first architecture



Data movement free from disk bottlenecks



- In-memory streaming of updates to all components
- In-memory cache
- Memory-only data buckets
- Memory-only indexes

In-memory streaming of updates to all components

In-memory (cached) access to data and indexes


Memory-only indexes

#3 Asynchronous approach to everything



- Persistence
- Intra-cluster Replication
- Inter-cluster Replication
- Global secondary Indexing updates
- Full-Text Search update
- Analytics service updates

#3 Asynchronous approach to everything



#3 Asynchronous approach to everything



Configurable consistency per request / query

Data Consistency

Data access is strongly consistent within cluster
Eventually consistent across clusters

Query Consistency

Specify level of consistency for queries

Thank you

