

# Architecture and Administration Basics

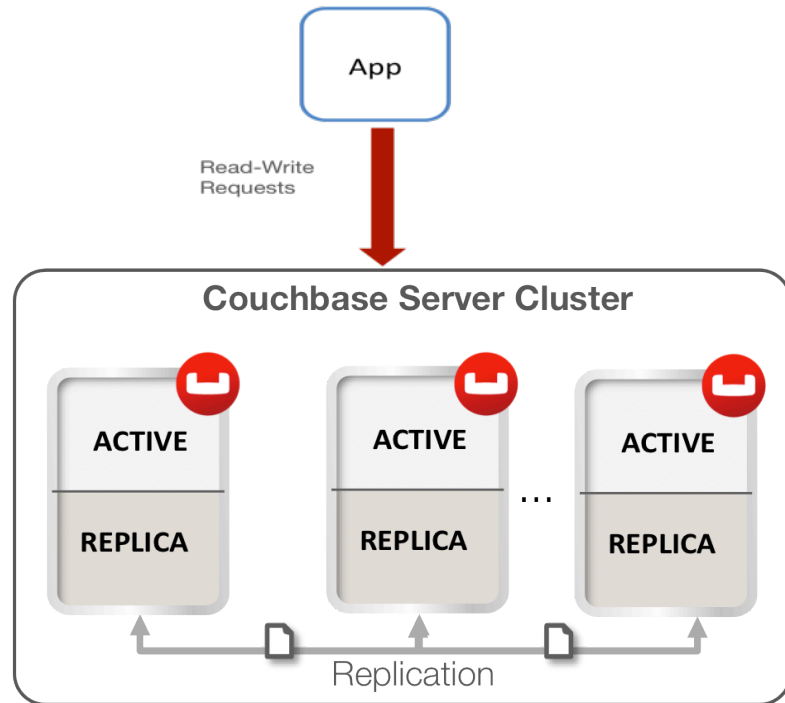
Workshop Day 1 - XDCR



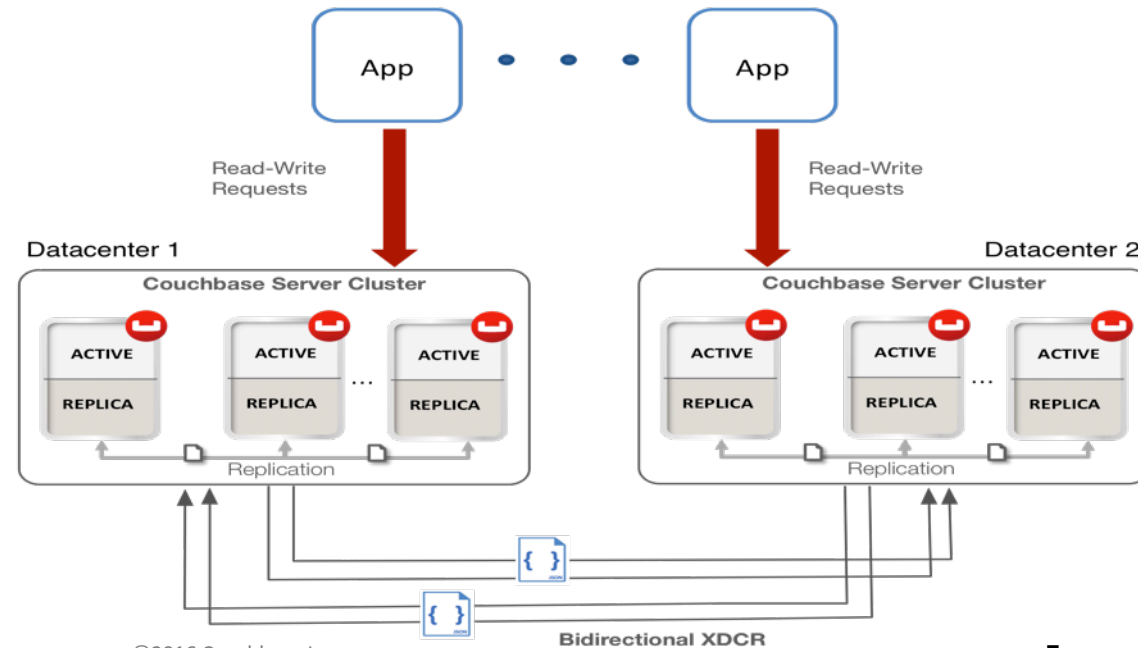
# 1

# Introduction

# Intra-Cluster vs. Inter-Cluster



VS.





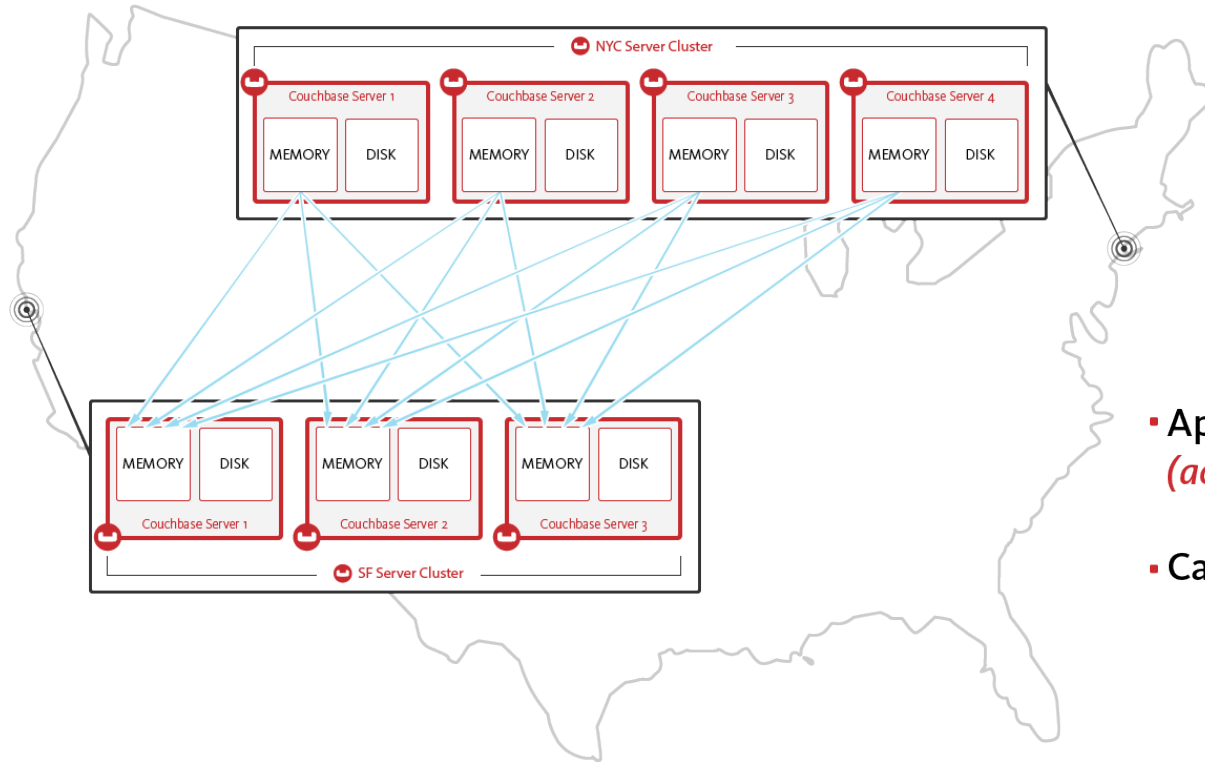
# Purposes

---

- Deliver high performing, async. data replication
- Provide disaster recovery and high availability across data centers
- Support data locality
- For load separation
- Support various topologies and replication schemes, including filtering
- Easy setup of development and test environments



# Example



- Applications can access both clusters (*active – active*)
- Can provide filtered replication




# 2

## Key Features



- **Continuous Replication**

- For existing and modified data
- Per bucket
- Asynchronous
- From memory
- Multiple data streams (configurable), shuffled across all shards to move data in parallel to the destination cluster
- Replication evenly load balanced across all servers in the destination cluster



- **Cluster Aware**

- Source and destinations can have different number of servers
- Takes topology update into account if E.G. a node of the destination cluster goes down



- **Automatic Resume**

- Push based replication
- Compares revisions before transfer
- Source tracks what destination last received via checkpoints
- Resume from last checkpoint

- **Conflict Resolution**


- Same 'winner' on both sides

- **Transfer and Security**

- HTTP (v1)
- XMEM (v2, uses Memcached protocol)
- Encrypted transfer (SSL)

- **Administrative Interface**

- Web-UI, REST, CLI









# 3

## How it works

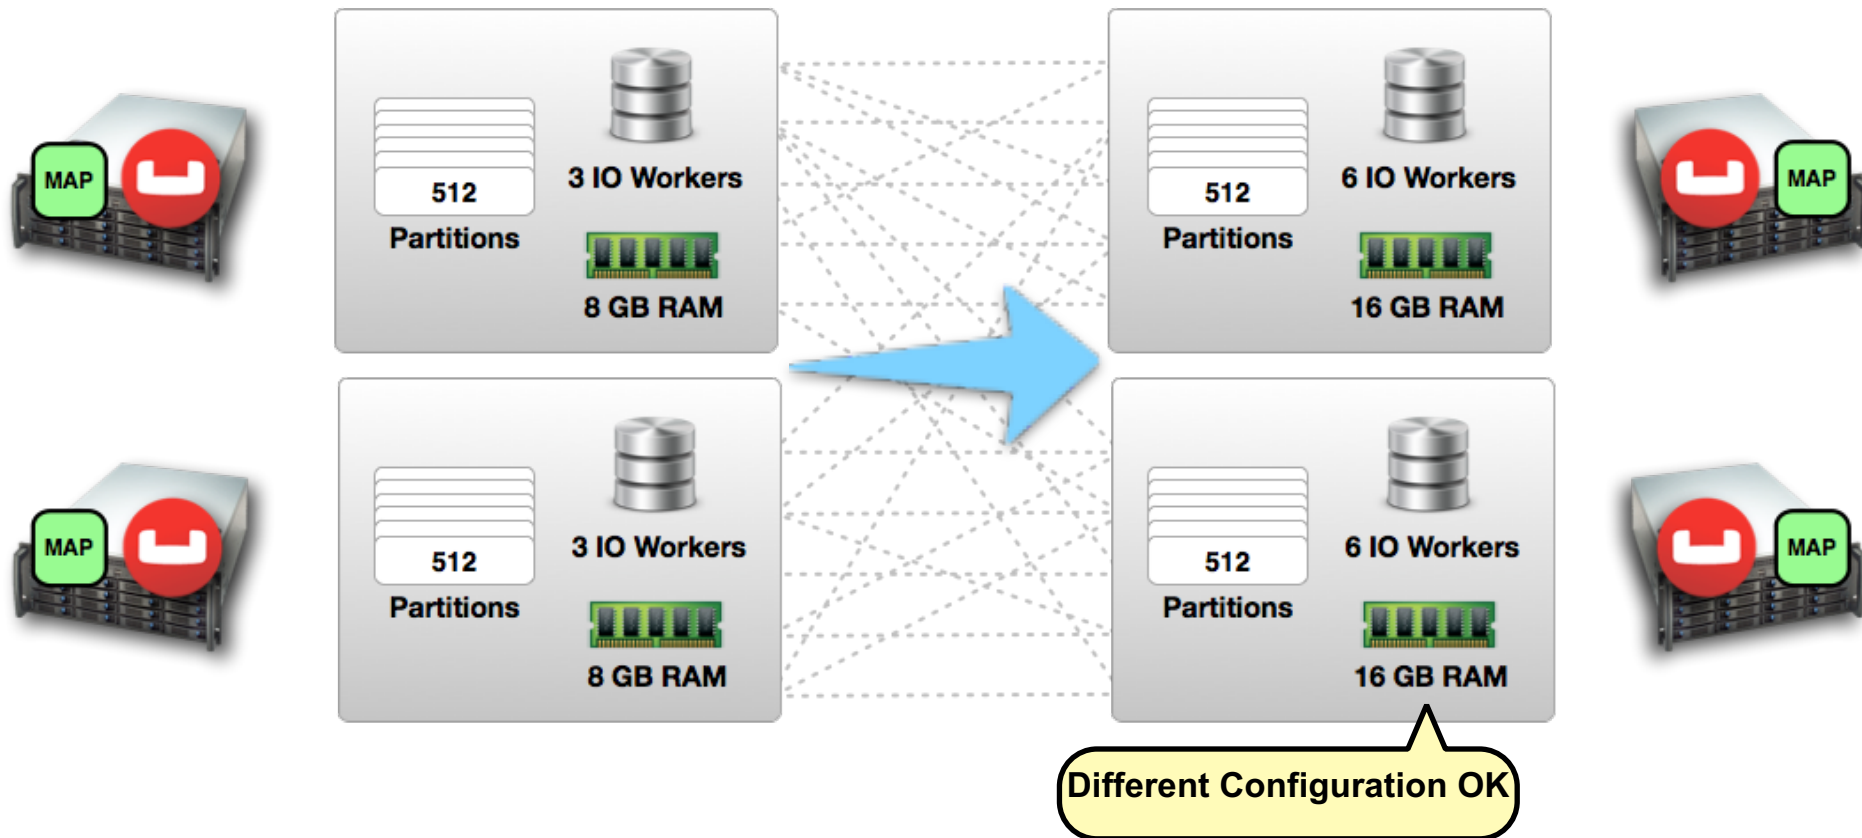
# XDCR after Write



# From Bucket to Bucket




## Follows Cluster Map



# Conflict resolution



- What happens when you write the same key in multiple clusters?




# Conflict resolution



- XDCR is eventually consistent; checks document metadata to resolve conflicts:

1. Numerical sequence (incremented on each mutation)
2. CAS value
3. Expiration (TTL) value

→ All clusters will pick the same “winner”





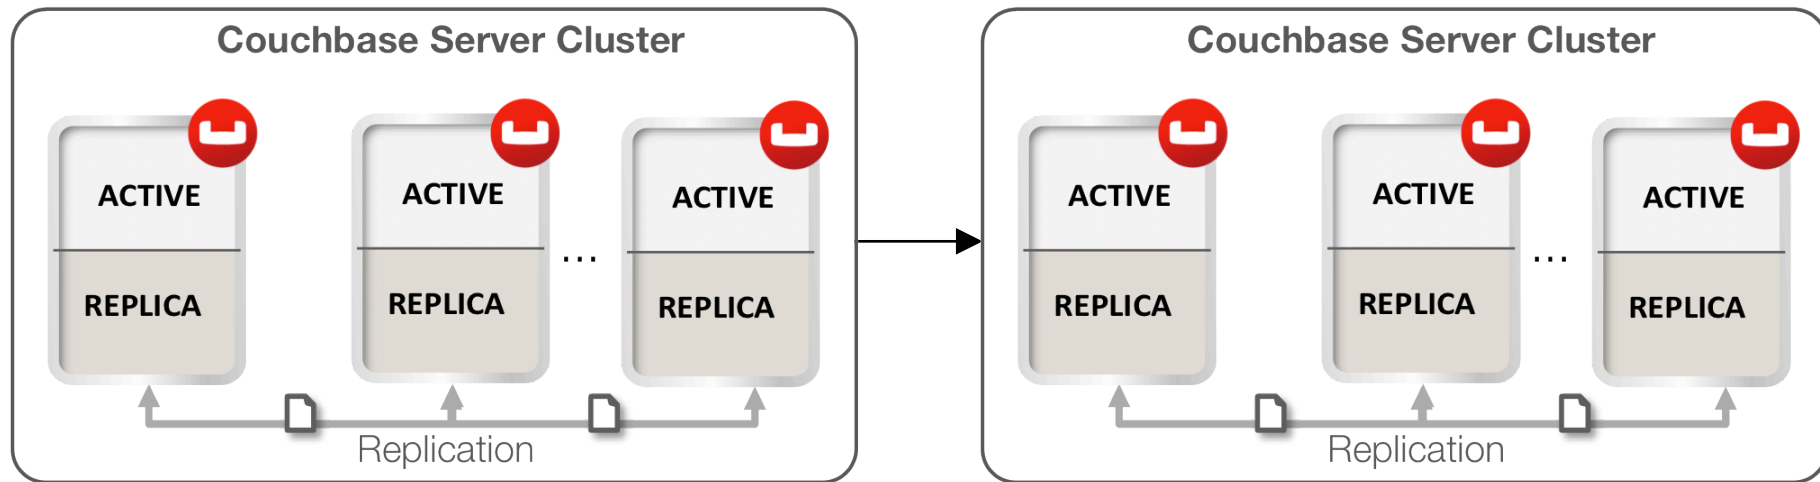
# 4

## Topologies & Use Cases

# Uni-Directional

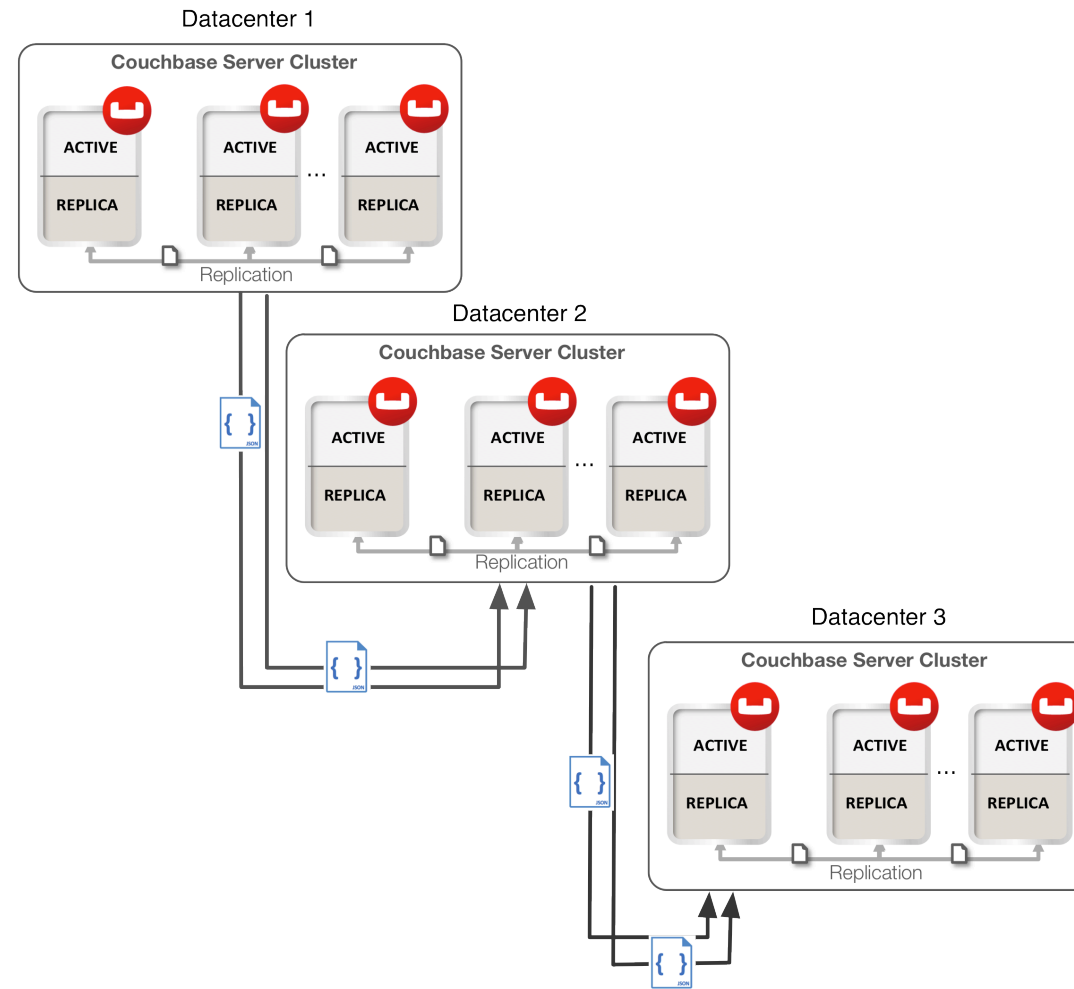


- Hot spare / Disaster recovery
- Development/testing copies
- Heavy reporting (since 4.0 via MDS)
- Integrate to Elasticsearch
- Integrate to custom consumer

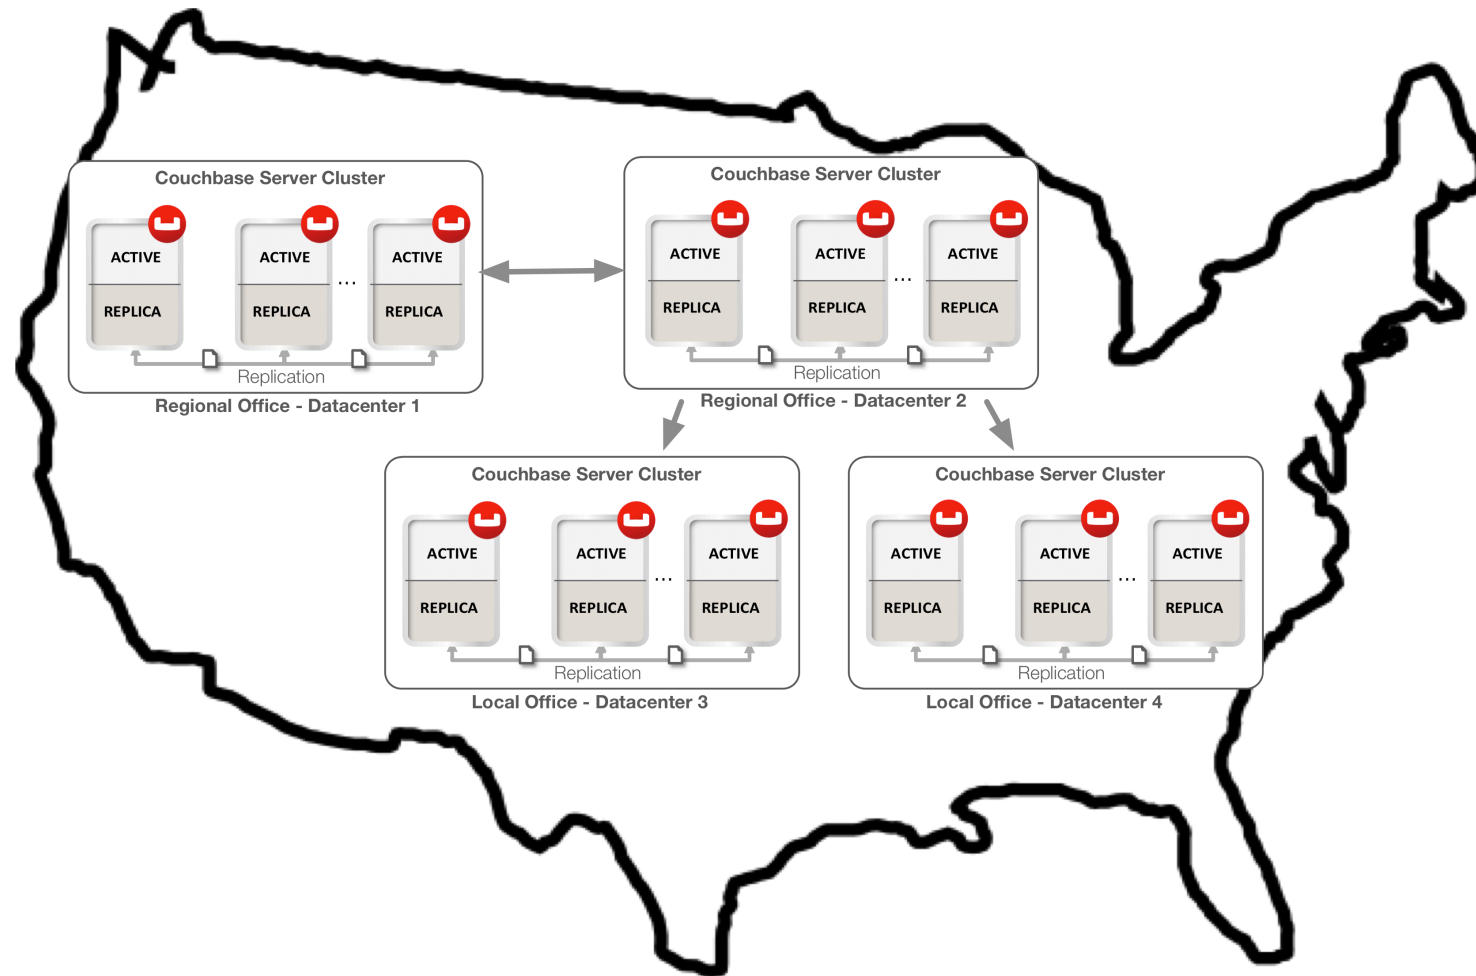




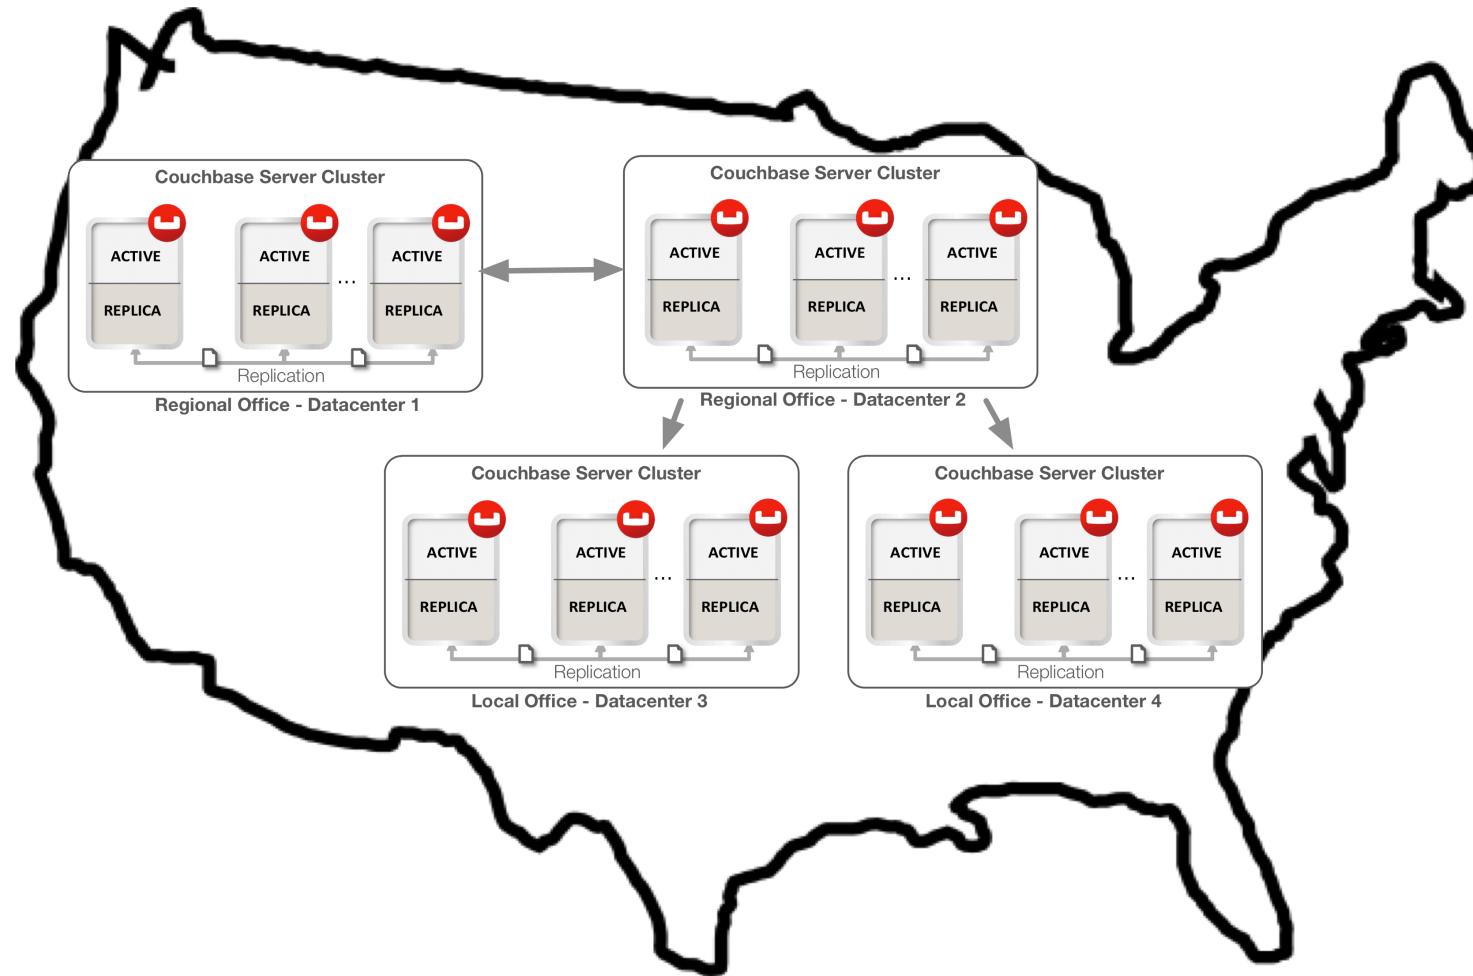
# Chain



# Data Aggregation



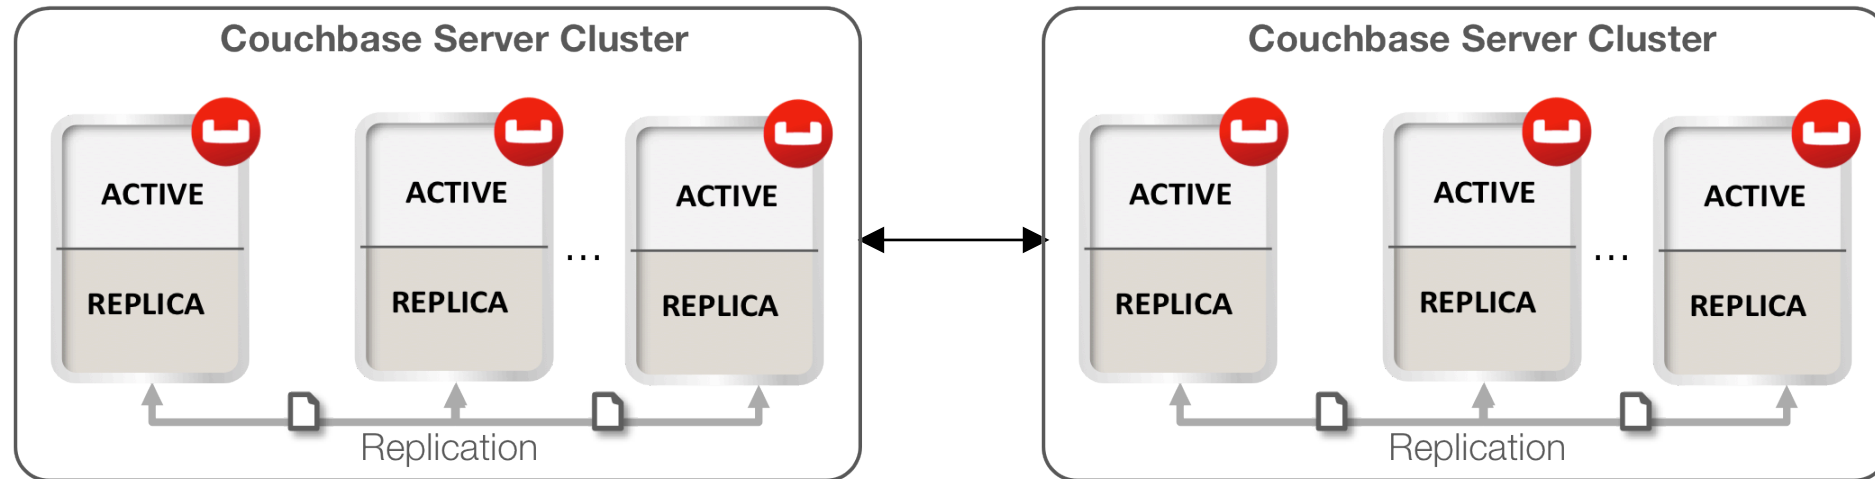
# (Filtered) Propagation





## Bi-Directional (aka Active-Active)

- Multiple active masters
- Disaster Recovery
- Data locality





- **Avoid updating the same document in multiple clusters with bi-directional XDCR**
  - Be sure to understand the conflict resolution rules
- **Best Practices**
  - Data Center stickiness
    - Keep users/transactions isolated to a DC
    - Only redirect to another DC in case of major outage
  - Use separate key spaces (e.g. DC prefix) to avoid conflicts on individual documents. Example:
    - dc1::user:a9838-s92-s00
    - dc2::user:293ba-293-922



# 5

## Tuning Parameters

# Advanced Settings



Parameter	Default	Description
Optimistic replication threshold	256	If the size of a document is higher than this threshold then XDCR will send a getMeta request (in batches) from the source cluster to the destination cluster in order to find out if the document needs to be sent over.
Source nozzles per node	2	Controls the parallelism
Target nozzles per node	2	Controls the parallelism
Checkpoint interval	1800	Time in seconds between checkpoints. This defines the amount of data which has to be resent in case of a communication failure.
Batch count	500	Controls the number of documents to be transferred in one batch.
Batch size (kB)	2048	Limits the size of a batch in KB.
Failure retry interval (s)	10	Time in seconds before XDCR retries to resume the replication after a failure.
Filter	None	The filter expression allows you to limit the data which will be sent over the wire by using a regular expression on the document key.

# Thank you

