

Architecture and Administration Basics

Workshop Day 2 - Json Data Modeling



1

JSON Document Design

JSON is all the rage these days



- JSON has (thankfully) replaced XML as the most common format used in APIs
- JSON is particularly useful for JavaScript apps since the data can be easily serialized/deserialized
- JSON has a number of supporting technologies such as JSONPath, JSON Schema and GeoJSON
- JSON does have several worthy competitors including BSON, YAML, Avro, Protocol Buffers and MessagePack
- JSON serialization libraries exist for most popular programming languages



- Single Root Attributes vs. "type"/"class" parameter
- Objects vs. Arrays
- Array Element Types
- Timestamp Formats
- Property Names
- Empty and Null Property Values
- JSON Schema

Single Root Attributes



- A choice between two styles:

```
{
  track: {
    artist: "Paul Lekakis",
    created: "2015-08-18T19:57:07",
    genre: "Hi-NRG",
    id: "3305311F4A0FAAFEABD001D324906748B18FB24A",
    mp3: https://goo.gl/KgKoR7,
    ratings: [
      {
        created: "2015-08-20T12:24:44",
        rating: 4,
        username: "sublimatingraga37014"
      },
      {
        created: "2015-08-21T09:23:57",
        rating: 4,
        username: "untillableshowings34122"
      },
      {
        created: "2015-08-21T13:53:34",
        rating: 3,
        username: "megacephalousfusty75226"
      }
    ],
    title: "My House",
    updated: "2015-08-18T19:57:07"
  }
}
```

```
{
  artist: "Paul Lekakis",
  created: "2015-08-18T19:57:07",
  genre: "Hi-NRG",
  id: "3305311F4A0FAAFEABD001D324906748B18FB24A",
  mp3: https://goo.gl/KgKoR7,
  ratings: [
    {
      created: "2015-08-20T12:24:44",
      rating: 4,
      username: "sublimatingraga37014"
    },
    {
      created: "2015-08-21T09:23:57",
      rating: 4,
      username: "untillableshowings34122"
    },
    {
      created: "2015-08-21T13:53:34",
      rating: 3,
      username: "megacephalousfusty75226"
    }
  ],
  title: "My House",
  type: "track",
  updated: "2015-08-18T19:57:07"
}
```

Objects vs. Arrays



- Two different ways to represent attributes:

```
{
  ▾ userprofile: {
    ▸ address: {...},
    created: "2015-01-28T13:50:56",
    dateOfBirth: "1986-06-09",
    email: "andy.bowman@games.com",
    ▸ favoriteGenres: [...],
    firstName: "Andy",
    gender: "male",
    lastName: "Bowman",
    ▾ phones: {
      cell: "212-771-1834"
    },
    ▸ picture: {...},
    pwd: "636f6c6f7261646f",
    status: "active",
    title: "Mr",
    updated: "2015-08-25T10:29:16",
    username: "copilotmarks61569"
  }
}
```

```
{
  ▾ userprofile: {
    ▸ address: {...},
    created: "2015-01-28T13:50:56",
    dateOfBirth: "1986-06-09",
    email: "andy.bowman@games.com",
    ▸ favoriteGenres: [...],
    firstName: "Andy",
    gender: "male",
    lastName: "Bowman",
    ▾ phones: [
      ▾ {
        number: "212-771-1834",
        type: "cell"
      }
    ],
    ▸ picture: {...},
    pwd: "636f6c6f7261646f",
    status: "active",
    title: "Mr",
    updated: "2015-08-25T10:29:16",
    username: "copilotmarks61569"
  }
}
```

Array Element Types



- Array elements can be simple types, objects or arrays:

```
{
  ▼ playlist: {
    created: "2014-12-04T03:36:18",
    id: "003c6f65-641a-4c9a-8e5e-41c947086cae",
    name: "Eclectic Summer Mix",
    owner: "copilotmarks61569",
    ▼ tracks: [
      "9FFAF88C1C3550245A19CE3BD91D3DC0BE616778",
      "3305311F4A0FAAFEABD001D324906748B18FB24A",
      "0EB4939F29669774A19B276E60F0E7B47E7EAF58"
    ],
    updated: "2015-09-11T10:39:40",
    visibility: "PUBLIC"
  }
}
```

Array of strings

```
{
  ▼ playlist: {
    created: "2014-12-04T03:36:18",
    id: "003c6f65-641a-4c9a-8e5e-41c947086cae",
    name: "Eclectic Summer Mix",
    owner: "copilotmarks61569",
    ▼ tracks: [
      ▼ {
        id: "9FFAF88C1C3550245A19CE3BD91D3DC0BE616778"
      },
      ▼ {
        id: "3305311F4A0FAAFEABD001D324906748B18FB24A"
      },
      ▼ {
        id: "0EB4939F29669774A19B276E60F0E7B47E7EAF58"
      }
    ],
    updated: "2015-09-11T10:39:40",
    visibility: "PUBLIC"
  }
}
```

Array of objects

Timestamp Formats



- When storing timestamps, you have at least 3 options:

```
{
  country: {
    countryCode: "US",
    gdp: 53548,
    name: "United States of America",
    population: 325296592,
    region: "Americas",
    region-number: 21,
    sub-region: "Northern America",
    updated: "2010-07-15T15:34:27"
  }
}
```

String (ISO 8601)

```
{
  country: {
    countryCode: "US",
    gdp: 53548,
    name: "United States of America",
    population: 325296592,
    region: "Americas",
    region-number: 21,
    sub-region: "Northern America",
    updated: 1279208067000
  }
}
```

Number (Unix style)

```
{
  country: {
    countryCode: "US",
    gdp: 53548,
    name: "United States of America",
    population: 325296592,
    region: "Americas",
    region-number: 21,
    sub-region: "Northern America",
    updated: [
      2010,
      7,
      15,
      15,
      34,
      27
    ]
  }
}
```

Array of time components

- Make timestamp values relative to UTC



- Choose meaningful property names
- Be consistent in naming properties
 - e.g. `country_code` vs. `countryCode` (preferred)
- Array types should have plural property names
- All other property names should be singular
- Avoid (if possible) reserved words in your database system and programming language(s)
 - e.g. ``user`` // Reserved word in Couchbase Server
- Avoid (if possible) special characters such as hypens
 - e.g. ``region-number`` // Contains a hyphen

Empty and Null Property Values



- Keep in mind that JSON supports optional properties
- If a property has a null value, consider dropping it from the JSON, unless there's a good reason not to
- N1QL makes it easy to test for missing or null property values

```
SELECT * FROM couchmusic1 WHERE userprofile.address IS NULL;
```

```
SELECT * FROM couchmusic1 WHERE userprofile.gender IS MISSING;
```

- Be sure your application code handles the case where a property value is missing



- Couchbase Server pays absolutely no attention to the shape of your JSON documents so long as they are well-formed
- There are times when it is useful to validate that a JSON document conforms to some expected shape
- JSON Schema is a JSON-based format for defining the structure of JSON data
- There are implementations for most popular programming languages
- Learn more here: <http://json-schema.org>

Example of JSON Schema



```
{
  id: "http://couchmusic.org/schema/couchmusic2-country.json",
  $schema: http://json-schema.org/draft-04/schema#,
  type: "object",
  properties: {
    countryCode: {
      type: "string",
      minLength: 2,
      maxLength: 2
    },
    gdp: {
      type: "integer",
      minimum: 0
    },
    name: {
      type: "string"
    },
    population: {
      type: "number",
      minimum: 0
    },
    region-number: {
      type: "integer",
      minimum: 0
    },
    type: {
      enum: [
        "country"
      ]
    },
    updated: {
      type: "string",
      format: "date-time"
    }
  },
  required: [
    "countryCode",
    "gdp",
    "name",
    "population",
    "region-number",
    "updated"
  ],
  additionalProperties: false
}
```

Example of JSON Schema – Type Specification



```
{
  id: "http://couchmusic.org/schema/couchmusic2-country.json",
  $schema: http://json-schema.org/draft-04/schema#,
  type: "object",
  properties: {
    countryCode: {
      type: "string",
      minLength: 2,
      maxLength: 2
    },
    gdp: {
      type: "integer",
      minimum: 0
    },
    name: {
      type: "string"
    },
    population: {
      type: "number",
      minimum: 0
    },
    region-number: {
      type: "integer",
      minimum: 0
    },
    type: {
      enum: [
        "country"
      ]
    },
    updated: {
      type: "string",
      format: "date-time"
    }
  },
  required: [
    "countryCode",
    "gdp",
    "name",
    "population",
    "region-number",
    "updated"
  ],
  additionalProperties: false
}
```

Available type specifications include:

- array
- boolean
- integer
- number
- object
- string
- enum

Example of JSON Schema – Type Specific Validation



```
{
  id: "http://couchmusic.org/schema/couchmusic2-country.json",
  $schema: http://json-schema.org/draft-04/schema#,
  type: "object",
  properties: {
    countryCode: {
      type: "string",
      minLength: 2,
      maxLength: 2
    },
    gdp: {
      type: "integer",
      minimum: 0
    },
    name: {
      type: "string"
    },
    population: {
      type: "number",
      minimum: 0
    },
    region-number: {
      type: "integer",
      minimum: 0
    },
    type: {
      enum: [
        "country"
      ]
    },
    updated: {
      type: "string",
      format: "date-time"
    }
  },
  required: [
    "countryCode",
    "gdp",
    "name",
    "population",
    "region-number",
    "updated"
  ],
  additionalProperties: false
}
```

Type specific validations include:

- minimum
- maximum
- minLength
- maxLength
- format
- pattern

Example of JSON Schema – Required Properties



```
{
  id: "http://couchmusic.org/schema/couchmusic2-country.json",
  $schema: http://json-schema.org/draft-04/schema#,
  type: "object",
  properties: {
    countryCode: {
      type: "string",
      minLength: 2,
      maxLength: 2
    },
    gdp: {
      type: "integer",
      minimum: 0
    },
    name: {
      type: "string"
    },
    population: {
      type: "number",
      minimum: 0
    },
    region-number: {
      type: "integer",
      minimum: 0
    },
    type: {
      enum: [
        "country"
      ]
    },
    updated: {
      type: "string",
      format: "date-time"
    }
  },
  required: [
    "countryCode",
    "gdp",
    "name",
    "population",
    "region-number",
    "updated"
  ],
  additionalProperties: false
}
```

Required properties can be specified for each object

Example of JSON Schema – Additional Properties



```
{
  id: "http://couchmusic.org/schema/couchmusic2-country.json",
  $schema: http://json-schema.org/draft-04/schema#,
  type: "object",
  properties: {
    countryCode: {
      type: "string",
      minLength: 2,
      maxLength: 2
    },
    gdp: {
      type: "integer",
      minimum: 0
    },
    name: {
      type: "string"
    },
    population: {
      type: "number",
      minimum: 0
    },
    region-number: {
      type: "integer",
      minimum: 0
    },
    type: {
      enum: [
        "country"
      ]
    },
    updated: {
      type: "string",
      format: "date-time"
    }
  },
  required: [
    "countryCode",
    "gdp",
    "name",
    "population",
    "region-number",
    "updated"
  ],
  additionalProperties: false
}
```

Additional properties can be disabled



2

Data Nesting



- As you know, relational database design promotes separating data using normalization, which doesn't scale
- For NoSQL systems, we often avoid normalization so that we can scale
- Nesting allows related objects to be organized into a hierarchical tree structure where you can have multiple levels of grouping
- Rule of thumb is to nest no more than 3 levels deep unless there is a very good reason to do so
- You will often want to include a timestamp in the nested data

Example #1 of Data Nesting



- Playlist with owner attribute containing username of corresponding userprofile

```
{
  ▾ playlist: {
    created: "2014-12-04T03:36:18",
    id: "003c6f65-641a-4c9a-8e5e-41c947086cae",
    name: "Eclectic Summer Mix",
    owner: "copilotmarks61569",
    ▸ tracks: [...],
    updated: "2015-09-11T10:39:40",
    visibility: "PUBLIC"
  }
}
```

```
{
  ▾ userprofile: {
    ▸ address: {...},
    created: "2015-01-28T13:50:56",
    dateOfBirth: "1986-06-09",
    email: "andy.bowman@games.com",
    ▸ favoriteGenres: [...],
    firstName: "Andy",
    gender: "male",
    lastName: "Bowman",
    ▸ phones: {...},
    ▸ picture: {...},
    pwd: "636f6c6f7261646f",
    status: "active",
    title: "Mr",
    updated: "2015-08-25T10:29:16",
    username: "copilotmarks61569"
  }
}
```

Example #1 of Data Nesting



- Playlist with owner attribute containing a subset of the corresponding userprofile

```
{
  ▾ playlist: {
    created: "2014-12-04T03:36:18",
    id: "003c6f65-641a-4c9a-8e5e-41c947086cae",
    name: "Eclectic Summer Mix",
    ▾ owner: {
      created: "2015-01-28T13:50:56",
      firstName: "Andy",
      lastName: "Bowman",
      title: "Mr",
      updated: "2015-08-25T10:29:16",
      username: "copilotmarks61569"
    },
    tracks: [...],
    updated: "2015-09-11T10:39:40",
    visibility: "PUBLIC"
  }
}
```

```
{
  ▾ userprofile: {
    ▸ address: {...},
    created: "2015-01-28T13:50:56",
    dateOfBirth: "1986-06-09",
    email: "andy.bowman@games.com",
    ▸ favoriteGenres: [...],
    firstName: "Andy",
    gender: "male",
    lastName: "Bowman",
    ▸ phones: {...},
    ▸ picture: {...},
    pwd: "636f6c6f7261646f",
    status: "active",
    title: "Mr",
    updated: "2015-08-25T10:29:16",
    username: "copilotmarks61569"
  }
}
```

* Note the inclusion of the **updated** attribute

Example #2 of Data Nesting



- Playlist with tracks attribute containing an array of track IDs

```
{
  ▾ playlist: {
    created: 1417685778000,
    id: "003c6f65-641a-4c9a-8e5e-41c947086cae",
    name: "Eclectic Summer Mix",
    ▶ owner: {...},
    ▾ tracks: [
      "9FFAF88C1C3550245A19CE3BD91D3DC0BE616778",
      "3305311F4A0FAAFEABD001D324906748B18FB24A",
      "0EB4939F29669774A19B276E60F0E7B47E7EAF58"
    ],
    updated: 1441985980000,
    visibility: "PUBLIC"
  }
}
```

Example #2 of Data Nesting



- Playlist with tracks attribute containing an array of track objects

```
{
  ▾ playlist: {
    created: 1417685778000,
    id: "003c6f65-641a-4c9a-8e5e-41c947086cae",
    name: "Eclectic Summer Mix",
    ▶ owner: {...},
    ▾ tracks: [
      ▾ {
        artist: "Gene Harris",
        genre: "Jazz Blues",
        id: "9FFAF88C1C3550245A19CE3BD91D3DC0BE616778",
        mp3: https://goo.gl/DEYx4X,
        title: "Battle Hymn of the Republic",
        updated: 1445167377000
      },
      ▶ {...},
      ▶ {...}
    ],
    updated: 1441985980000,
    visibility: "PUBLIC"
  }
}
```

* Note the inclusion of the **updated** attribute



3

Key Design

- A key formed of attributes that exist in the real world:
 - Phone numbers
 - Usernames
 - Social security numbers
 - Account numbers
 - SKU, UPC or QR codes
 - Device IDs
- Often the first choice for document keys
- Be careful when working with any personally identifiable information (PII), sensitive personal information (SPI) or protected health information (PHI)





- We often use surrogate keys when no obvious natural key exist
- They are not derived from application data
- They can be generated values
 - 3305311F4A0FAAFEABD001D324906748B18FB24A (SHA-1)
 - 003C6F65-641A-4CGA-8E5E-41C947086CAE (UUID)
- They can be sequential numbers (often implemented using the Counter feature of Couchbase Server)
- 456789, 456790, 456791, ...



- Common practice for users of Couchbase Server to follow patterns for formatting key values by using symbols such as single or double colons
- DocType::ID
 - userprofile::fredsmith79
 - playlist::003c6f65-641a-4c9a-8e5e-41c947086cae
- AppName::DocType::ID
 - couchmusic::userprofile::fredsmith79
- DocType::ParentID::ChildID
 - playlist::fredsmith79::003c6f65-641a-4c9a-8e5e-41c947086cae
- Supports easy document viewing in the Couchbase web console



- The purpose of the Lookup Key Pattern is to allow multiple ways to reach the same data, essentially a secondary index
 - For example, we want to lookup a Userprofile by their email address instead of their ID
- To accomplish this, we create another small document that refers to the Userprofile document we are interested in
- Implementing this pattern is straightforward, just create an additional document containing a single property that stores the key to the primary document
- With the introduction of N1QL, this pattern will be less commonly used

Example of Lookup Key Pattern




- Lookup document can be JsonDocument or StringDocument

userprofile::copilotmarks61569

```
{
  ▾ userprofile: {
    ▸ address: { ... },
    created: "2015-01-28T13:50:56",
    dateOfBirth: "1986-06-09",
    email: "andy.bowman@games.com",
    ▸ favoriteGenres: [ ... ],
    firstName: "Andy",
    gender: "male",
    lastName: "Bowman",
    ▸ phones: { ... },
    ▸ picture: { ... },
    pwd: "636f6c6f7261646f",
    status: "active",
    title: "Mr",
    updated: "2015-08-25T10:29:16",
    username: "copilotmarks61569"
  }
}
```

andy.bowman@games.com

```
{
  username: "copilotmarks61569"
}
```



andy.bowman@games.com

```
copilotmarks61569
```





4

Making Trade-offs

Dictionary

trade-off

noun | \ˈtrād-,ôf\

: a situation in which you must choose between or balance two things that are opposite or cannot be had at the same time

: something that you do not want but must accept in order to have something that you want

- Eric Brewer is famous for showing the trade-offs that are necessary when dealing with distributed systems
 - Consistency, availability and partition tolerance are all desirable properties but we must choose the ones that are most important for our use cases
- We must also make trade-offs in data modeling:
 - Document size
 - Atomicity
 - Complexity
 - Speed





- Couchbase Server supports documents up to 20 Mb
- Larger documents take more disk space, more time to transfer across the network and more time to serialize/deserialize
- If you are dealing with documents that are potentially large (greater than 1 Mb), you must test thoroughly to find out if speed of access is adequate as you scale. If not, you will need to break up the document into smaller ones.
- You may need to limit the number of dependent child objects you embed



- Atomicity in Couchbase Server is at the document level
- Couchbase Server does not support transactions
 - They can be simulated if you are willing to write and maintain additional code to implement them (generally not recommended)
- If you absolutely need changes to be atomic, they will have to be part of the same document
- The maximum document size for Couchbase Server may limit how much data you can store in a single document



- Complexity affects every area of software systems including data modeling
- We need to consider:

```
– 1 SELECT c.name, COUNT(*) AS playlist_count
   FROM couchmusic1.playlist p JOIN couchmusic1.userprofile u
     ON KEYS 'userprofile::' || p.owner JOIN couchmusic1.country c
     ON KEYS 'country::' || u.address.countryCode
  WHERE c.`region-number` = 154
     AND p.visibility = 'PUBLIC'
     AND u.status = 'active'
  GROUP BY c.name
  ORDER BY c.name;
```

- The complexity of code for updating multiple copies of the same data




- As it relates to data modeling, speed of access is critical
- When using N1QL to access data, keep in mind that query by document key is fastest and query by secondary index is usually much slower
- If implementing an interactive use case, you will want to avoid using JOINS
- You can use data duplication to improve the speed of accessing related data and thus trade improved speed for greater complexity and larger document size
- Keep in mind that Couchbase Views can be used when up to the second accuracy is not required



- All of the previous trade-offs are usually rolled into a single decision – whether to embed or refer
- When to embed:
 - Reads greatly outnumber writes
 - You're comfortable with the slim risk of inconsistent data across the multiple copies
 - You're optimizing for speed of access
- When to refer:
 - Consistency of the data is a priority
 - You want to ensure your cache is used efficiently
 - The embedded version would be too large or complex



- In this module, you have learned to:
 - Make full use of JSON capabilities
 - Use data nesting to minimize the need for JOINS
 - Establish key value patterns and use them consistently
 - Be clear about the trade-offs you are making, document your decisions and the assumptions they are based on



Thank you

